

```
In [6]: import pandas as pd  
  
df = pd.read_csv("C:/Users/Admin/Downloads/index_data.csv")  
print(df.head())  
print(df.info())
```

	Creditability	Account Balance	Duration of Credit (month)	\
0	1	1	18	
1	1	1	9	
2	1	2	12	
3	1	1	12	
4	1	1	12	

	Payment Status of Previous Credit	Purpose	Credit Amount	\
0	4	2	1049	
1	4	0	2799	
2	2	9	841	
3	4	0	2122	
4	4	0	2171	

	Value Savings/Stocks	Length of current employment	Instalment per cent	\
0	1	2	4	
1	1	3	2	
2	2	4	2	
3	1	3	3	
4	1	3	4	

	Sex & Marital Status	...	Duration in Current address	\
0	2	...	4	
1	3	...	2	
2	2	...	4	
3	3	...	2	
4	3	...	4	

	Most valuable available asset	Age (years)	Concurrent Credits	\
0	2	21	3	
1	1	36	3	
2	1	23	3	
3	1	39	3	
4	2	38	1	

	Type of apartment	No of Credits at this Bank	Occupation	\
0	1	1	3	
1	1	2	3	
2	1	1	2	
3	1	2	2	
4	2	2	2	

	No of dependents	Telephone	Foreign Worker	
0	1	1	1	
1	2	1	1	
2	1	1	1	
3	2	1	2	
4	1	1	2	

[5 rows x 21 columns]
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 21 columns):
Column

0 Creditability

#	Column	Non-Null Count	Dtype
0	Creditability	1000 non-null	int64

```
1 Account Balance           1000 non-null int64
2 Duration of Credit (month) 1000 non-null int64
3 Payment Status of Previous Credit 1000 non-null int64
4 Purpose                  1000 non-null int64
5 Credit Amount             1000 non-null int64
6 Value Savings/Stocks      1000 non-null int64
7 Length of current employment 1000 non-null int64
8 Instalment per cent       1000 non-null int64
9 Sex & Marital Status      1000 non-null int64
10 Guarantors                1000 non-null int64
11 Duration in Current address 1000 non-null int64
12 Most valuable available asset 1000 non-null int64
13 Age (years)               1000 non-null int64
14 Concurrent Credits        1000 non-null int64
15 Type of apartment         1000 non-null int64
16 No of Credits at this Bank 1000 non-null int64
17 Occupation                 1000 non-null int64
18 No of dependents          1000 non-null int64
19 Telephone                  1000 non-null int64
20 Foreign Worker             1000 non-null int64
dtypes: int64(21)
memory usage: 164.2 KB
None
```

```
In [10]: print(df.describe(include='all'))
```

	Creditability	Account Balance	Duration of Credit (month)	\
count	1000.000000	1000.000000	1000.000000	
mean	0.700000	2.577000	20.903000	
std	0.458487	1.257638	12.058814	
min	0.000000	1.000000	4.000000	
25%	0.000000	1.000000	12.000000	
50%	1.000000	2.000000	18.000000	
75%	1.000000	4.000000	24.000000	
max	1.000000	4.000000	72.000000	

	Payment Status of Previous Credit	Purpose	Credit Amount	\
count	1000.000000	1000.000000	1000.000000	
mean	2.545000	2.828000	3271.24800	
std	1.08312	2.744439	2822.75176	
min	0.00000	0.000000	250.00000	
25%	2.00000	1.000000	1365.50000	
50%	2.00000	2.000000	2319.50000	
75%	4.00000	3.000000	3972.25000	
max	4.00000	10.000000	18424.00000	

	Value Savings/Stocks	Length of current employment	\
count	1000.000000	1000.000000	
mean	2.105000	3.384000	
std	1.580023	1.208306	
min	1.000000	1.000000	
25%	1.000000	3.000000	
50%	1.000000	3.000000	
75%	3.000000	5.000000	
max	5.000000	5.000000	

	Instalment per cent	Sex & Marital Status	...	\
count	1000.000000	1000.000000	...	
mean	2.973000	2.68200	...	
std	1.118715	0.70808	...	
min	1.000000	1.00000	...	
25%	2.000000	2.00000	...	
50%	3.000000	3.00000	...	
75%	4.000000	3.00000	...	
max	4.000000	4.00000	...	

	Duration in Current address	Most valuable available asset	\
count	1000.000000	1000.000000	
mean	2.845000	2.358000	
std	1.103718	1.050209	
min	1.000000	1.000000	
25%	2.000000	1.000000	
50%	3.000000	2.000000	
75%	4.000000	3.000000	
max	4.000000	4.000000	

	Age (years)	Concurrent Credits	Type of apartment	\
count	1000.000000	1000.000000	1000.000000	
mean	35.54200	2.675000	1.928000	
std	11.35267	0.705601	0.530186	
min	19.00000	1.000000	1.000000	
25%	27.00000	3.000000	2.000000	

50%	33.00000	3.000000	2.000000
75%	42.00000	3.000000	2.000000
max	75.00000	3.000000	3.000000

	No of Credits at this Bank	Occupation	No of dependents	Telephone	\
count	1000.00000	1000.00000	1000.00000	1000.00000	
mean	1.407000	2.904000	1.155000	1.404000	
std	0.577654	0.653614	0.362086	0.490943	
min	1.000000	1.000000	1.000000	1.000000	
25%	1.000000	3.000000	1.000000	1.000000	
50%	1.000000	3.000000	1.000000	1.000000	
75%	2.000000	3.000000	1.000000	2.000000	
max	4.000000	4.000000	2.000000	2.000000	

	Foreign Worker
count	1000.00000
mean	1.037000
std	0.188856
min	1.000000
25%	1.000000
50%	1.000000
75%	1.000000
max	2.000000

[8 rows x 21 columns]

In [11]: `print(df.isnull().sum())`

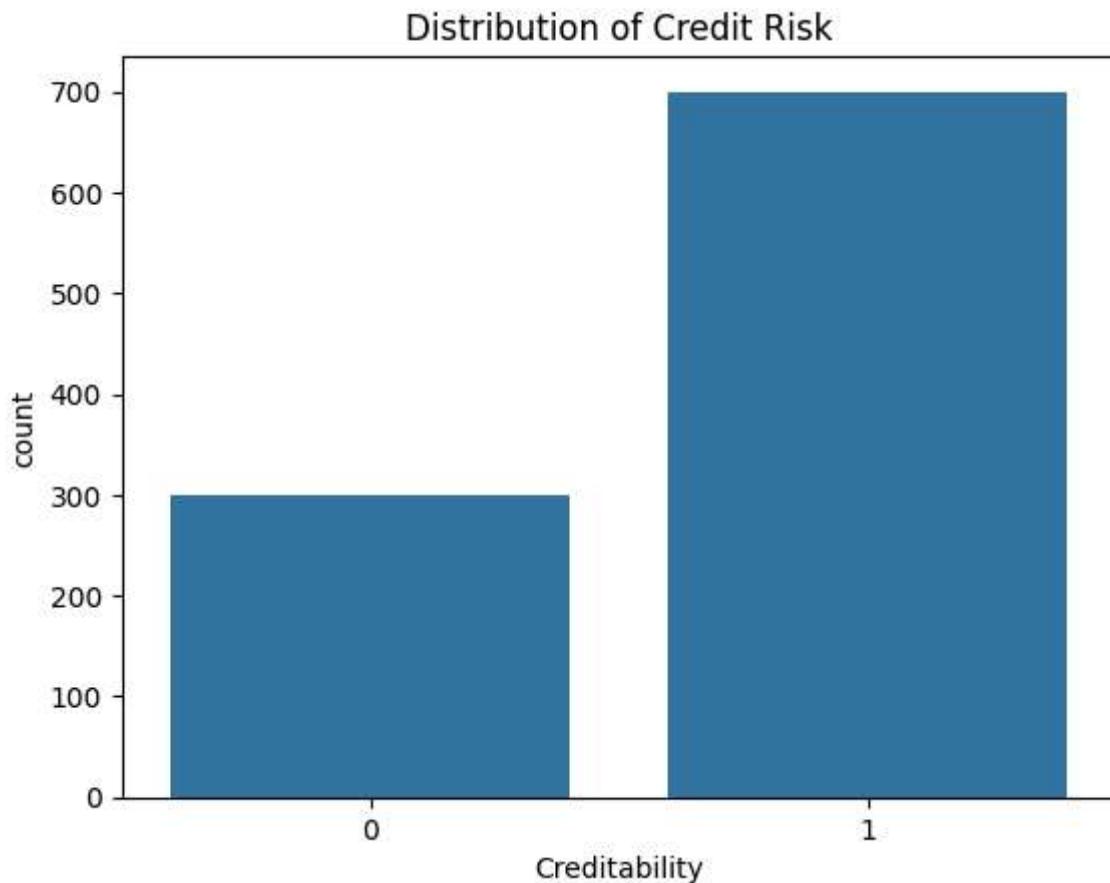
Creditability	0
Account Balance	0
Duration of Credit (month)	0
Payment Status of Previous Credit	0
Purpose	0
Credit Amount	0
Value Savings/Stocks	0
Length of current employment	0
Instalment per cent	0
Sex & Marital Status	0
Guarantors	0
Duration in Current address	0
Most valuable available asset	0
Age (years)	0
Concurrent Credits	0
Type of apartment	0
No of Credits at this Bank	0
Occupation	0
No of dependents	0
Telephone	0
Foreign Worker	0

`dtype: int64`

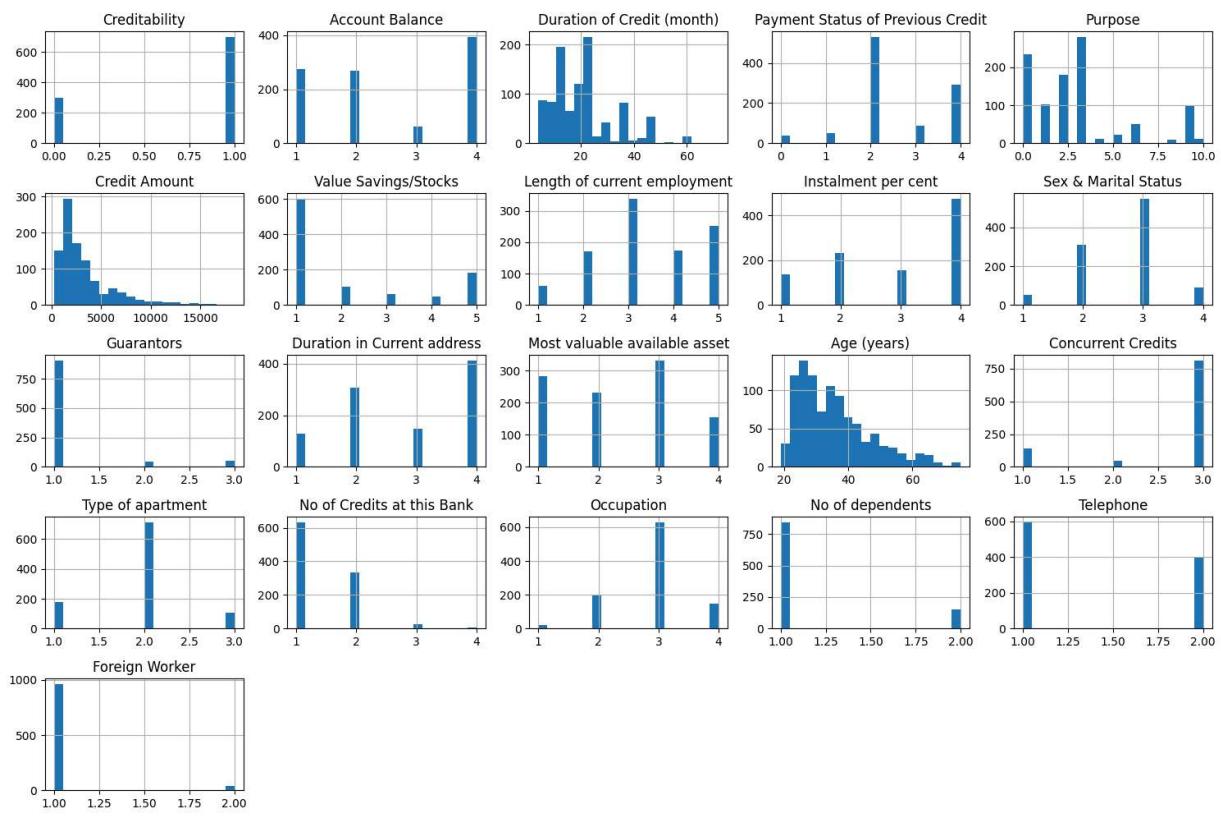
In [13]: `import seaborn as sns
import matplotlib.pyplot as plt`

`sns.countplot(x='Creditability', data=df)`

```
plt.title("Distribution of Credit Risk")
plt.show()
```



```
In [14]: df.hist(figsize=(15,10), bins=20)
plt.tight_layout()
plt.show()
```

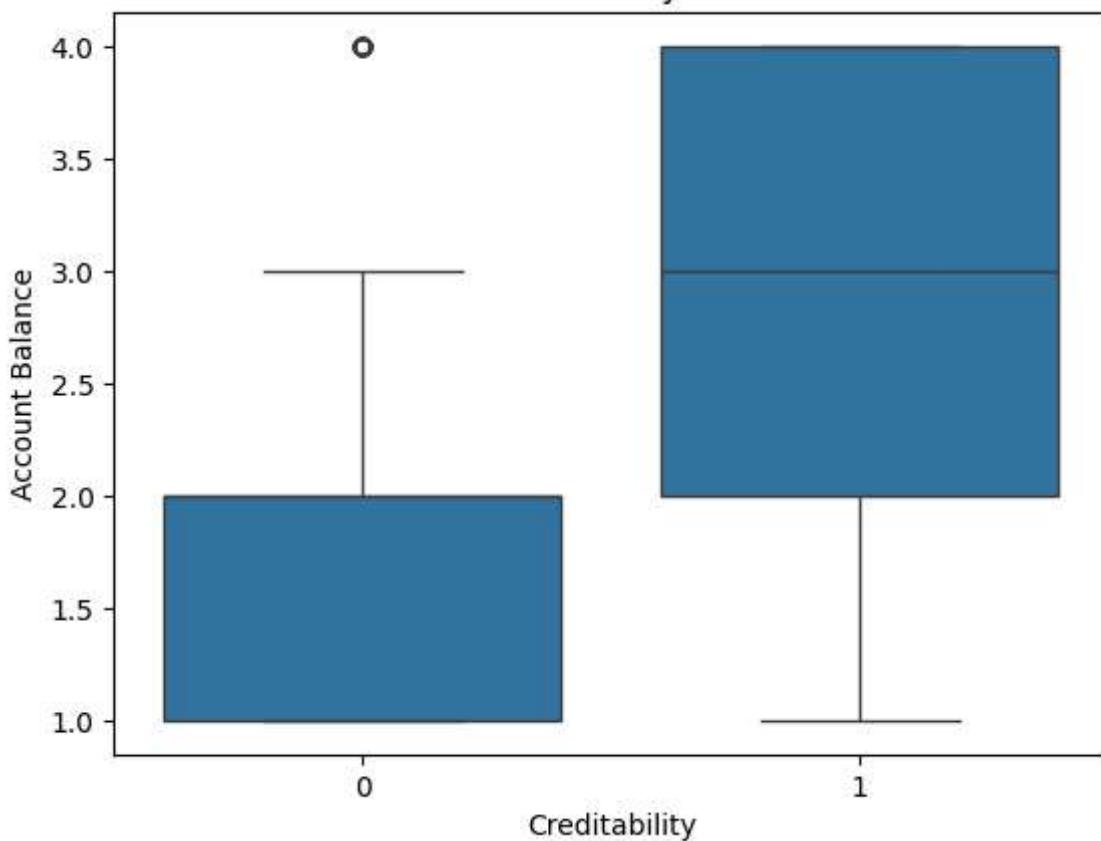


```
In [16]: for col in df.select_dtypes(include='object').columns:
    sns.countplot(y=col, data=df)
    plt.title(f"Distribution of {col}")
    plt.show()
```

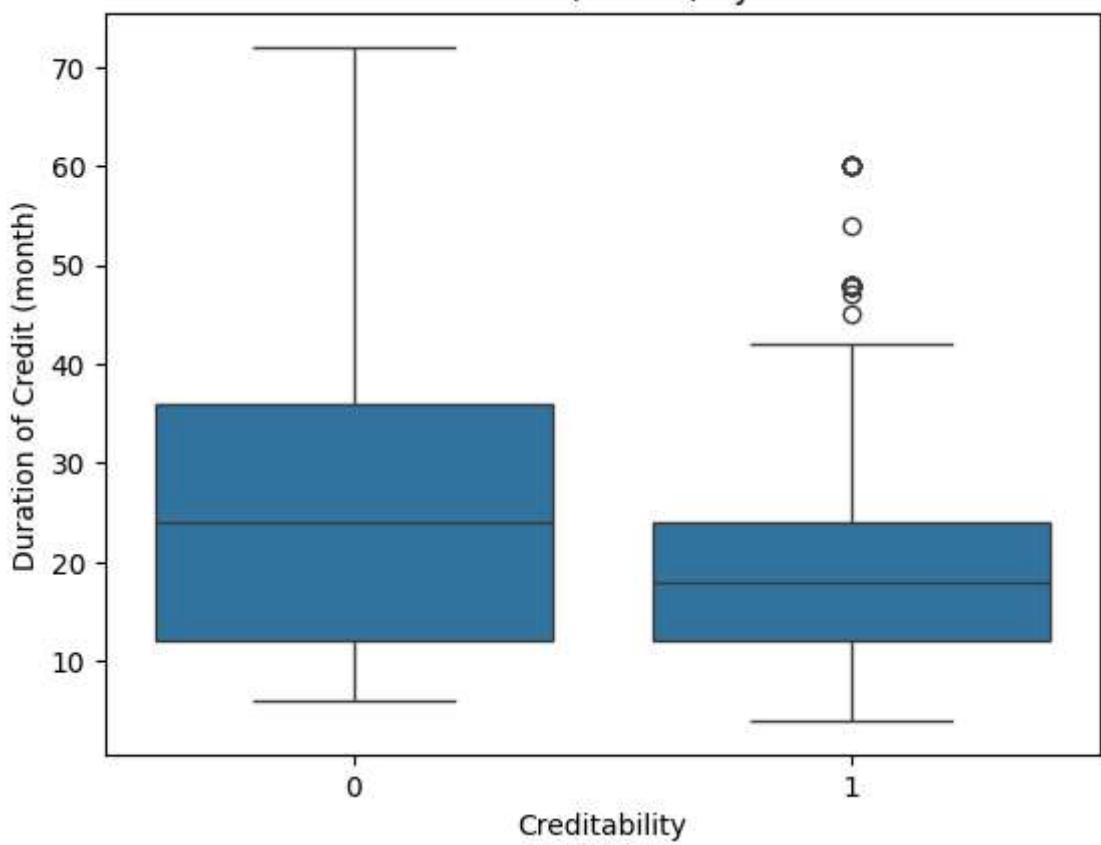
```
In [17]: import seaborn as sns

for col in df.select_dtypes(include='number').columns:
    if col != 'Creditability':
        sns.boxplot(x='Creditability', y=col, data=df)
        plt.title(f"{col} by Credit Risk")
        plt.show()
```

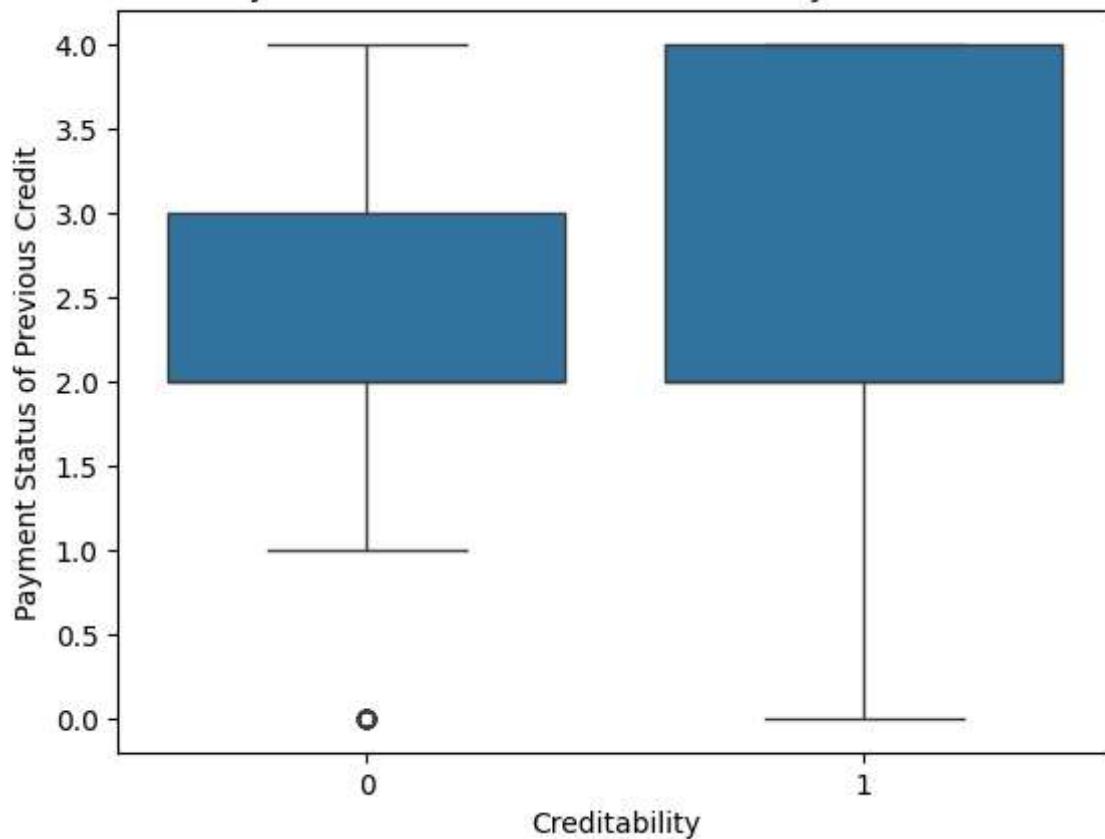
Account Balance by Credit Risk



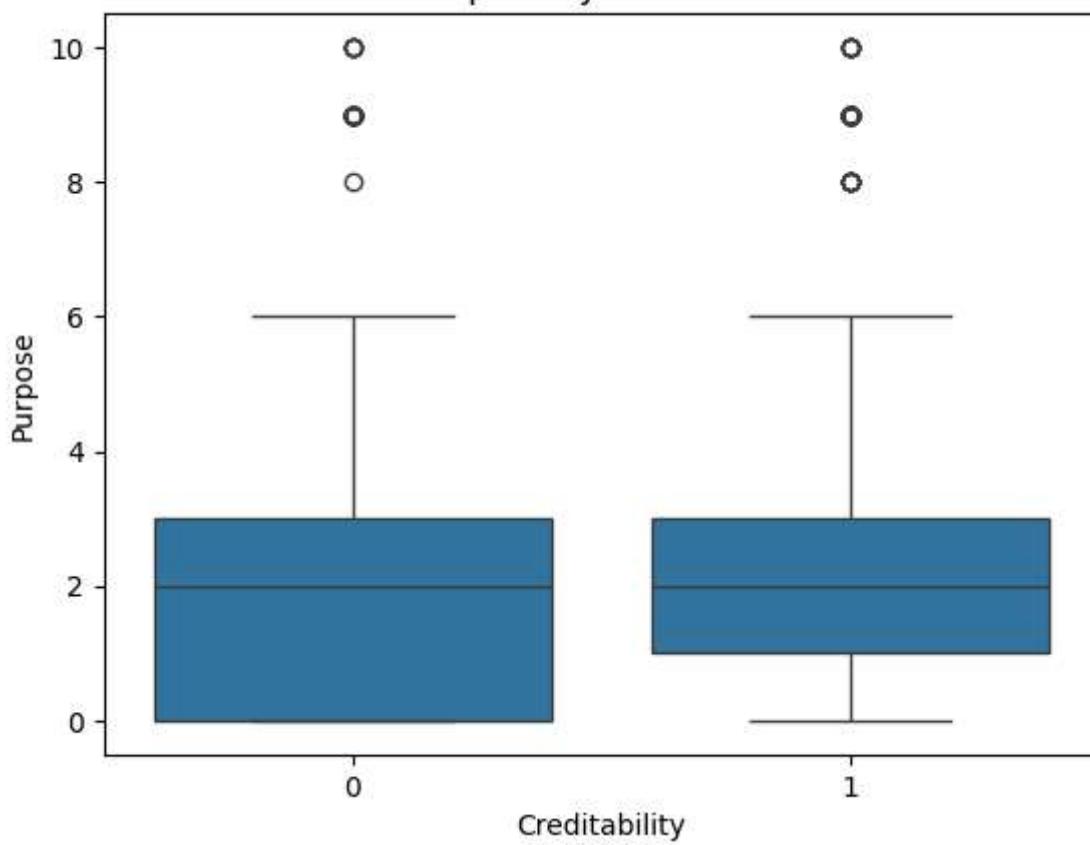
Duration of Credit (month) by Credit Risk



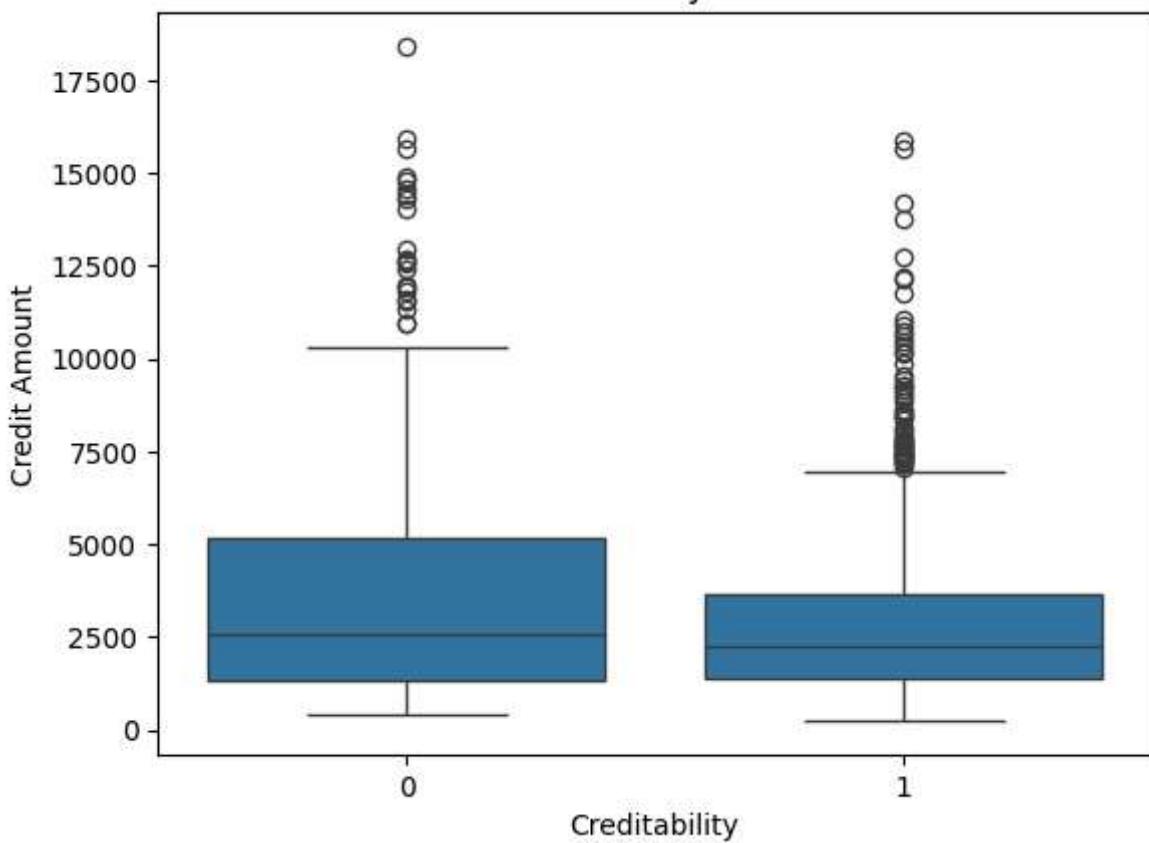
Payment Status of Previous Credit by Credit Risk



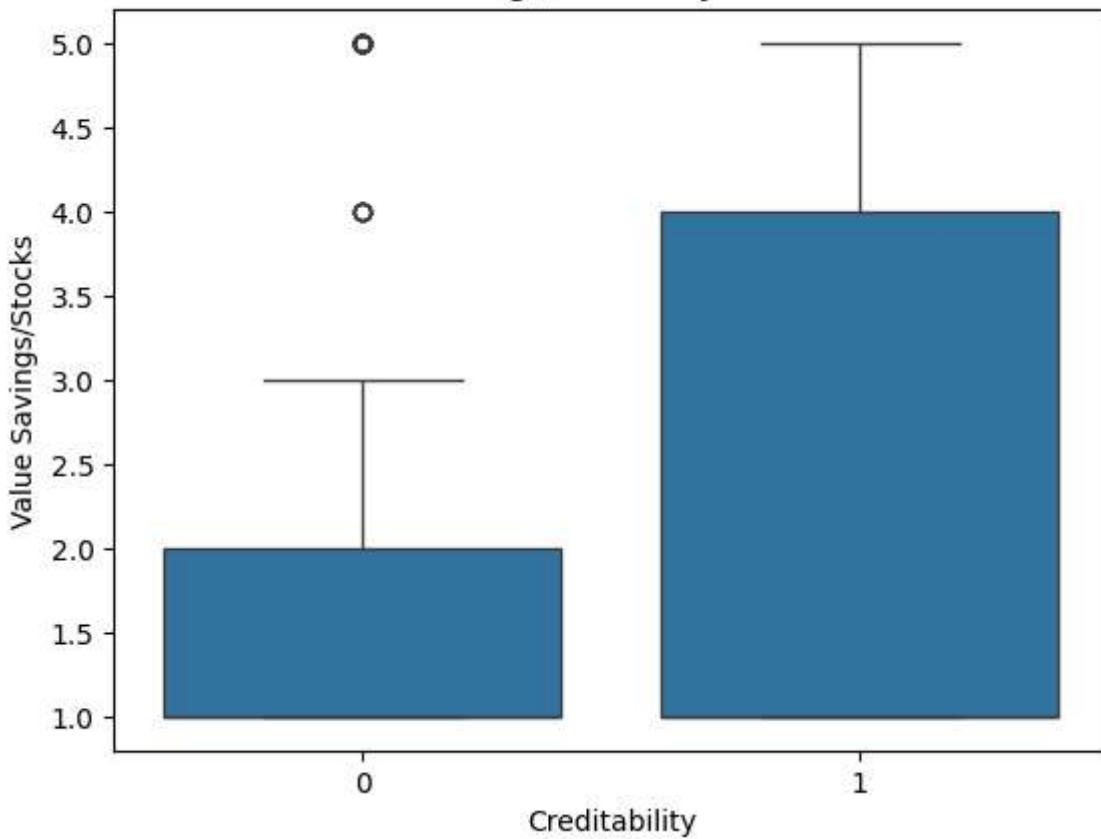
Purpose by Credit Risk



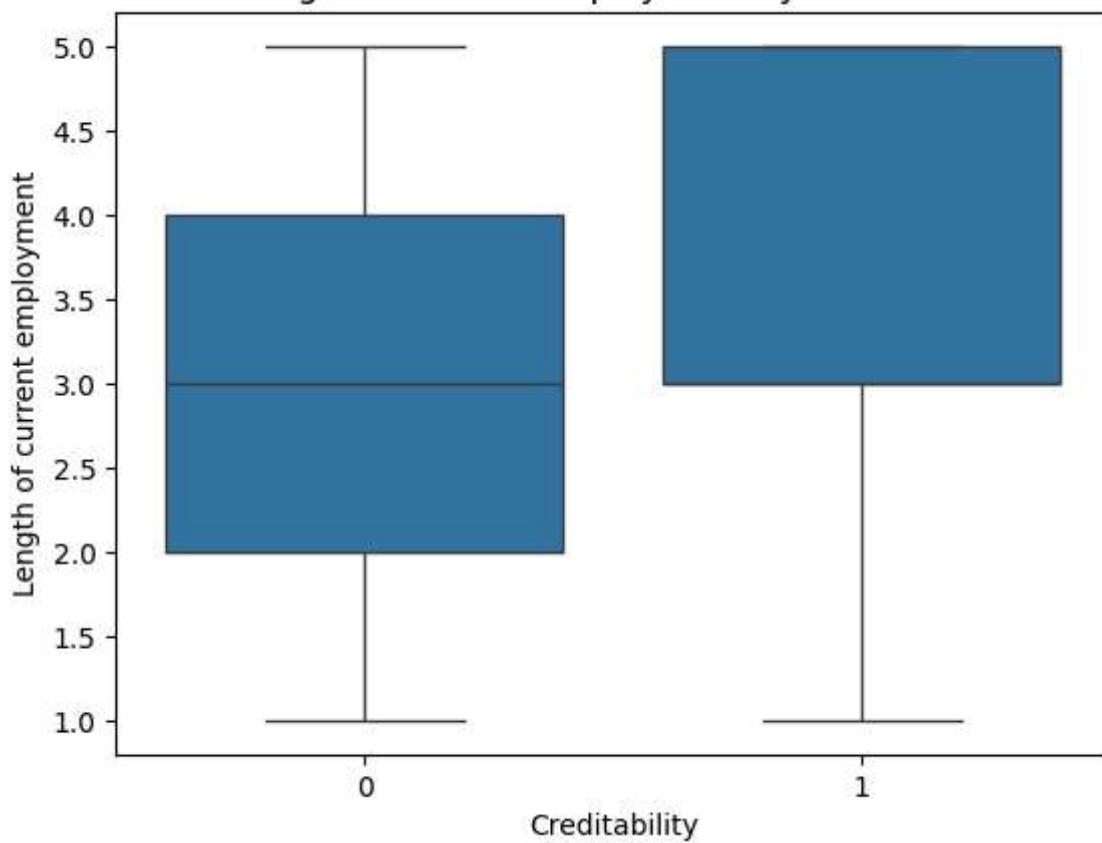
Credit Amount by Credit Risk



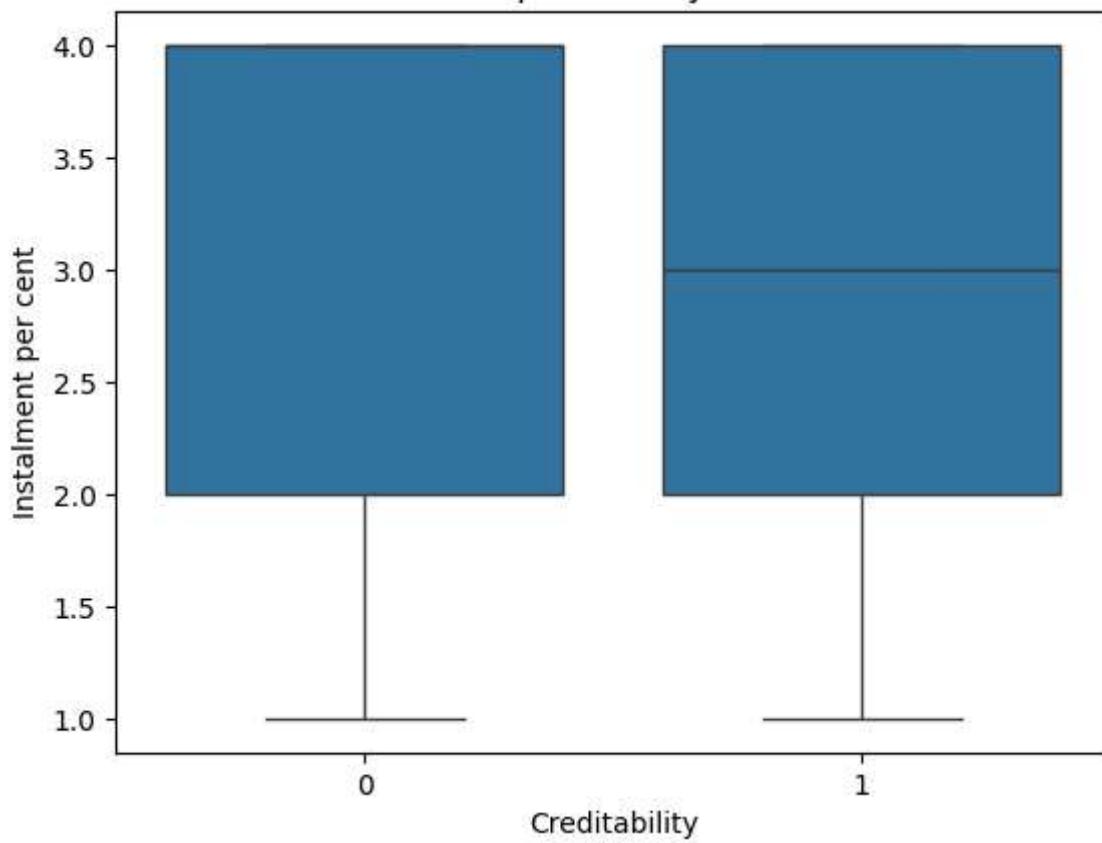
Value Savings/Stocks by Credit Risk



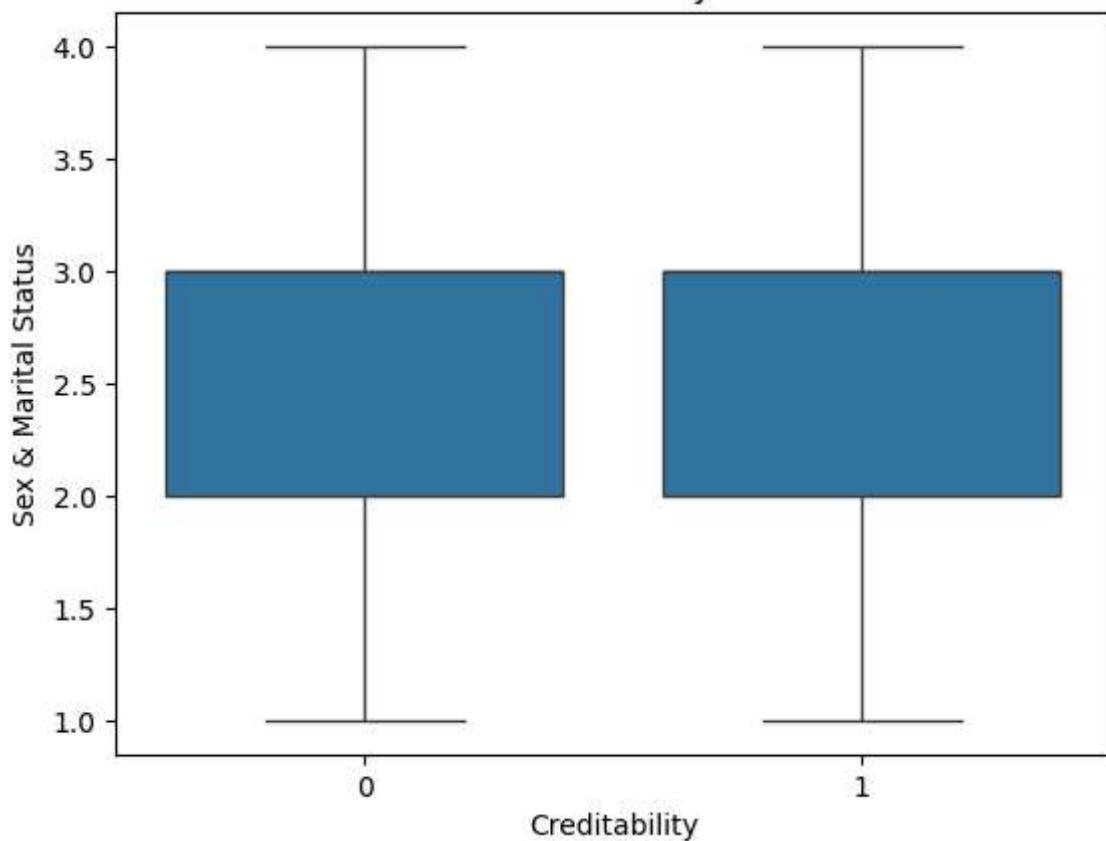
Length of current employment by Credit Risk



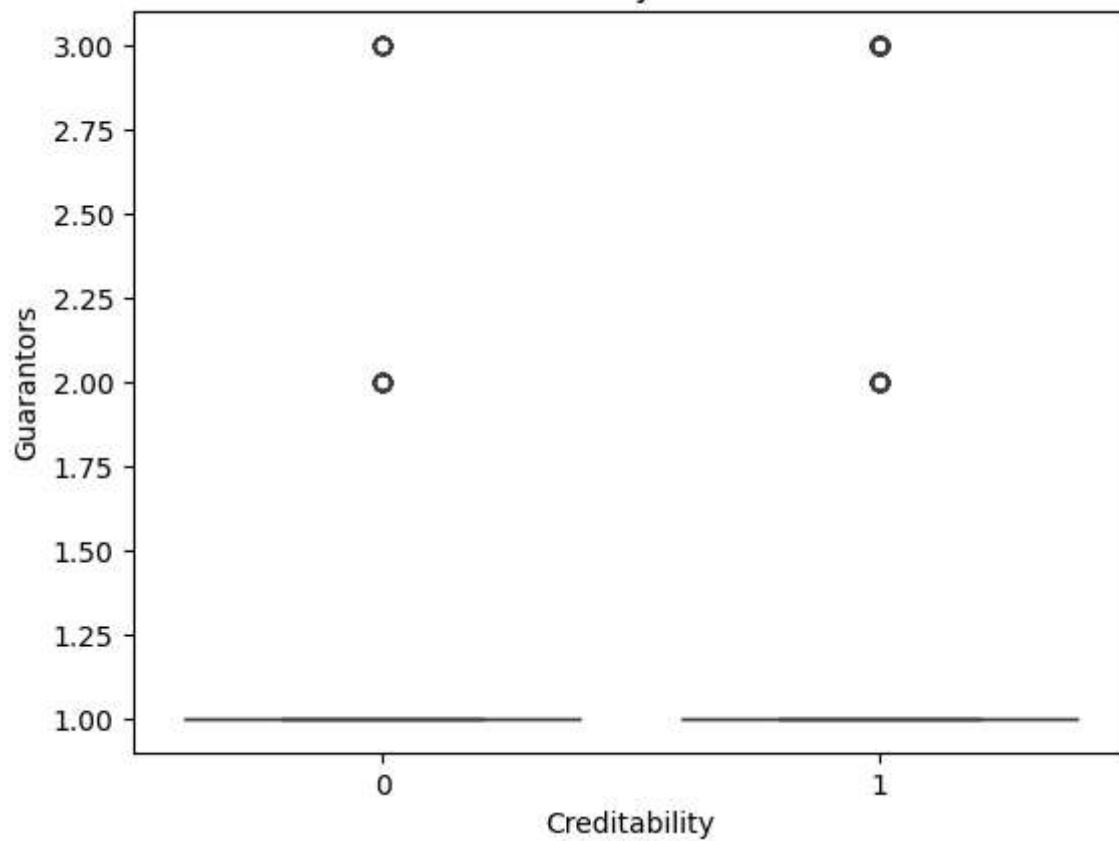
Instalment per cent by Credit Risk

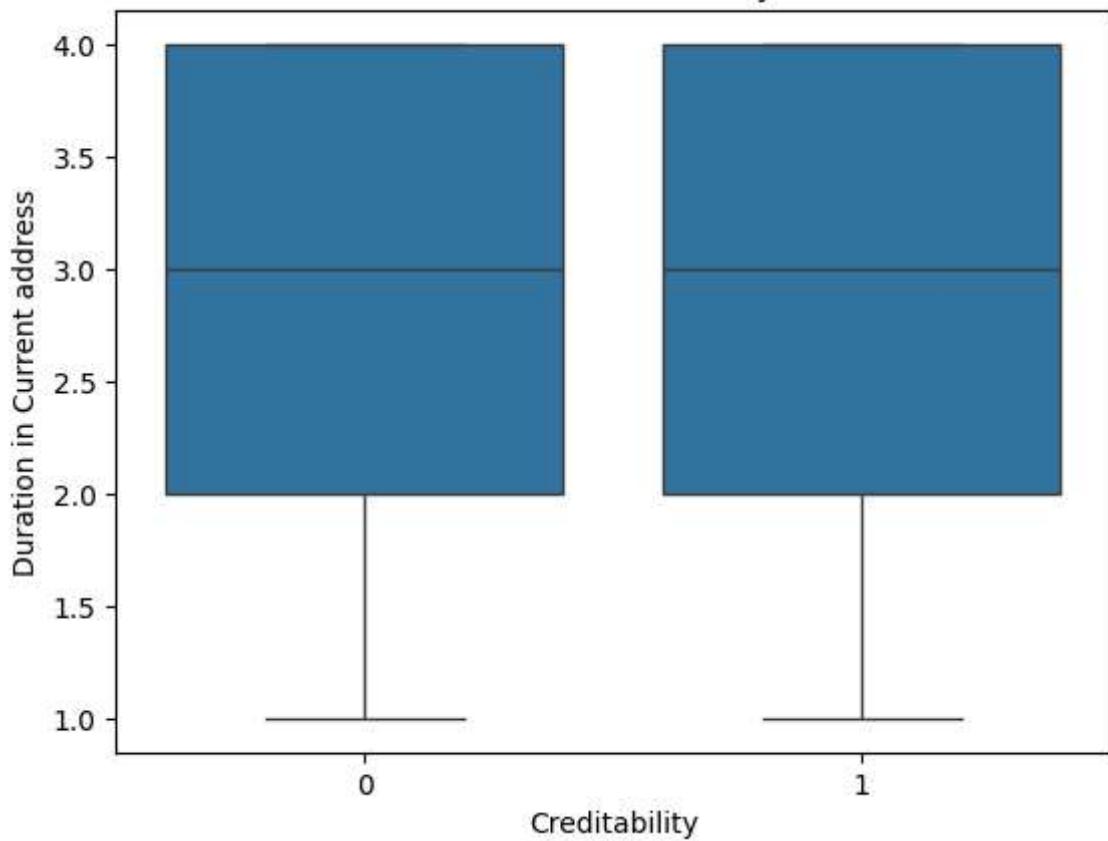
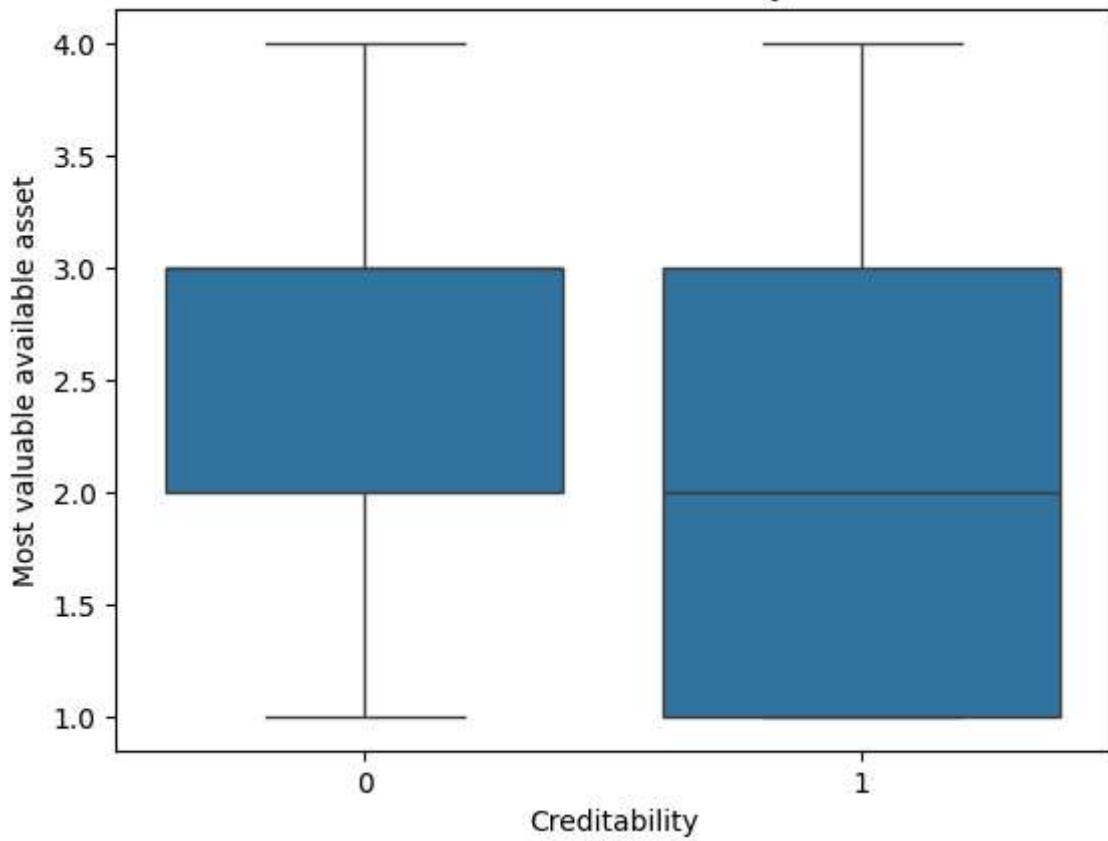


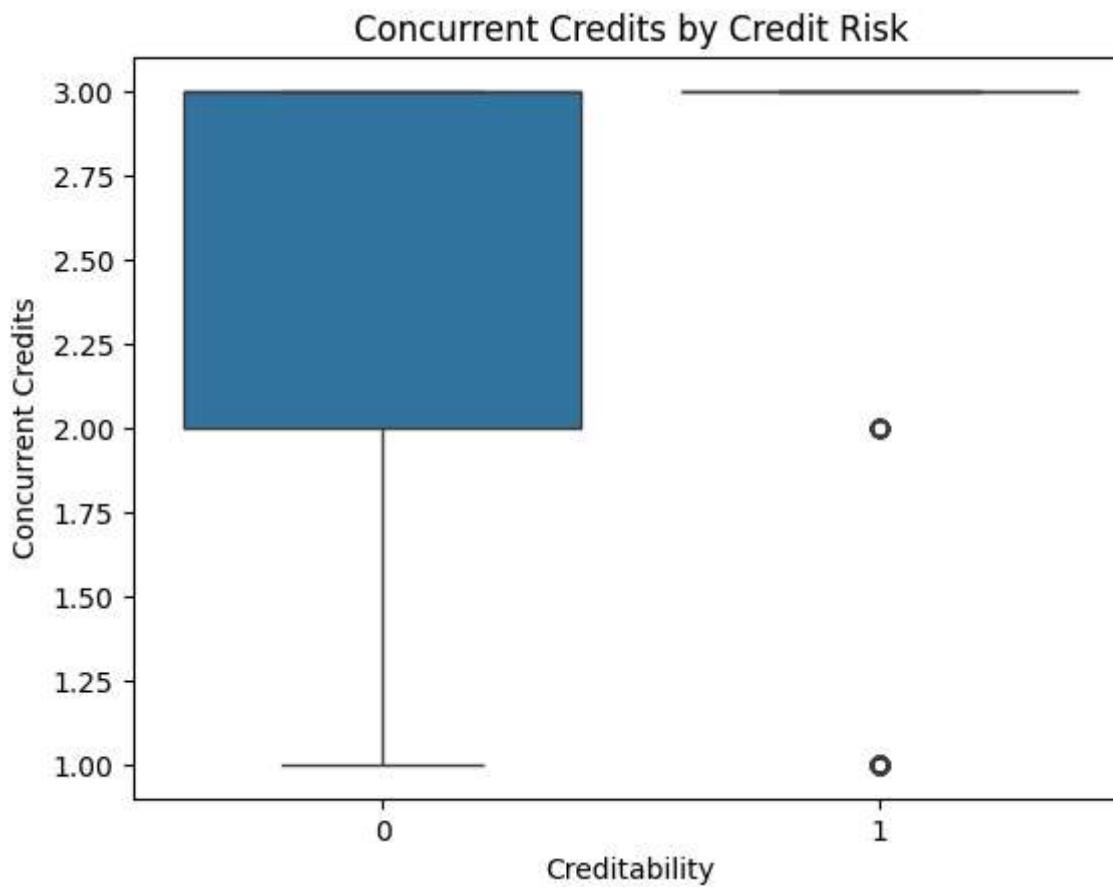
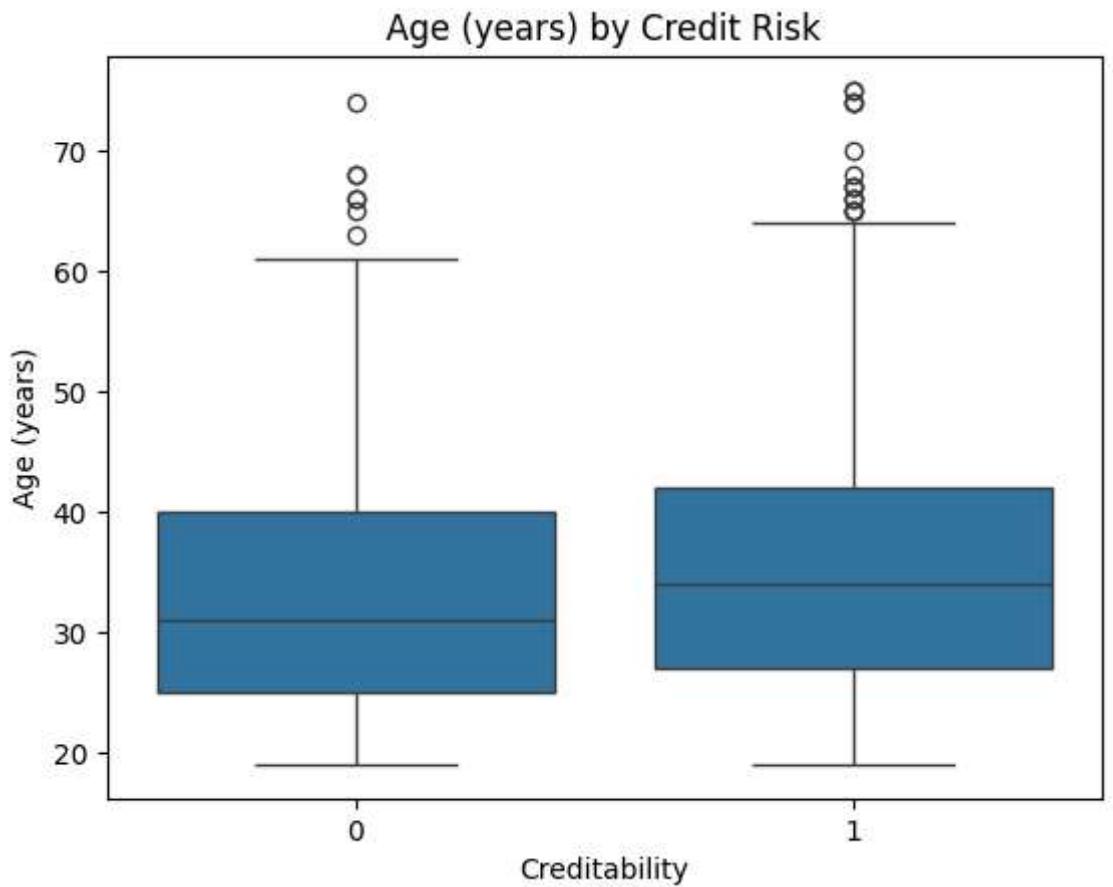
Sex & Marital Status by Credit Risk



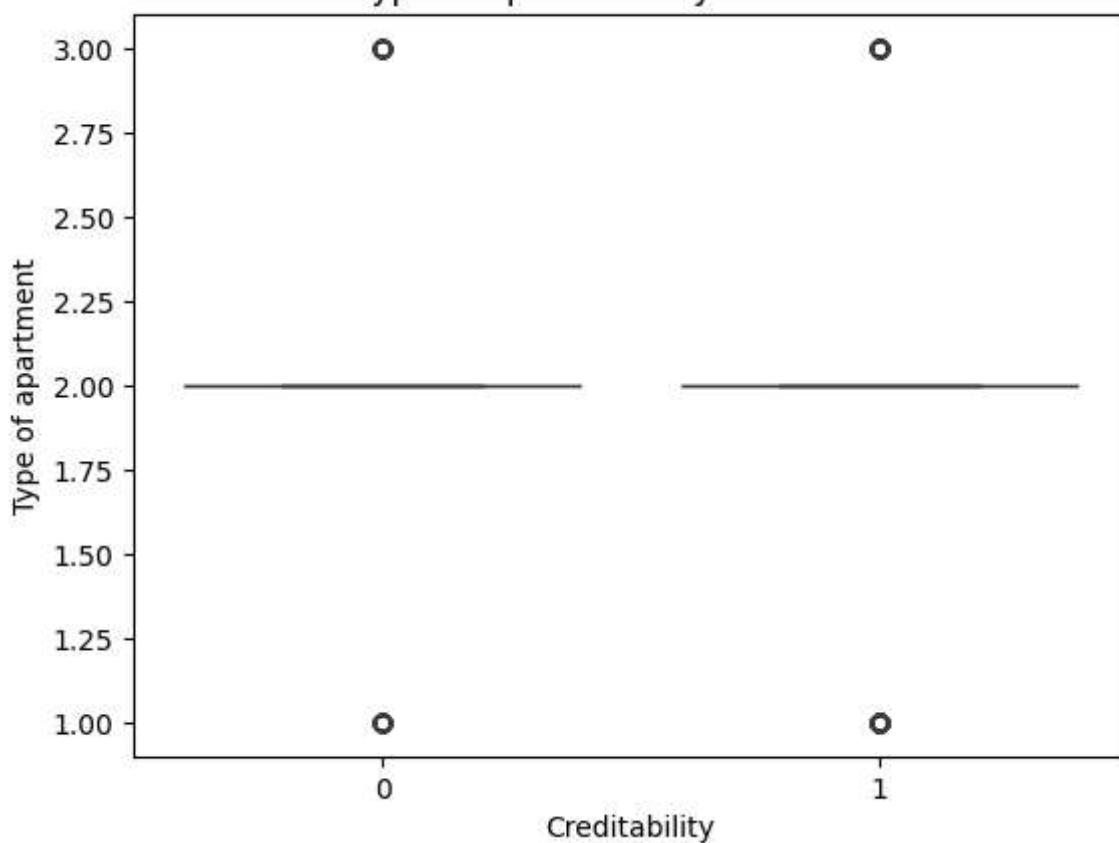
Guarantors by Credit Risk



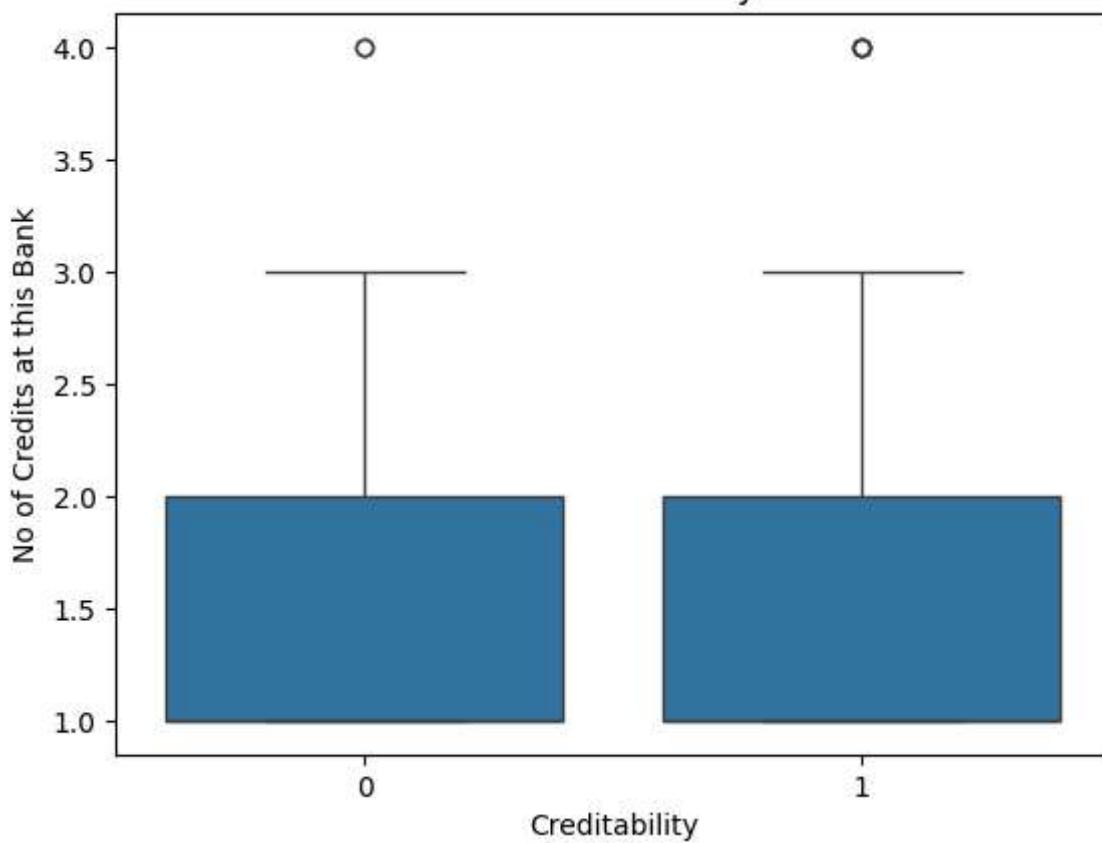
Duration in Current address by Credit Risk**Most valuable available asset by Credit Risk**

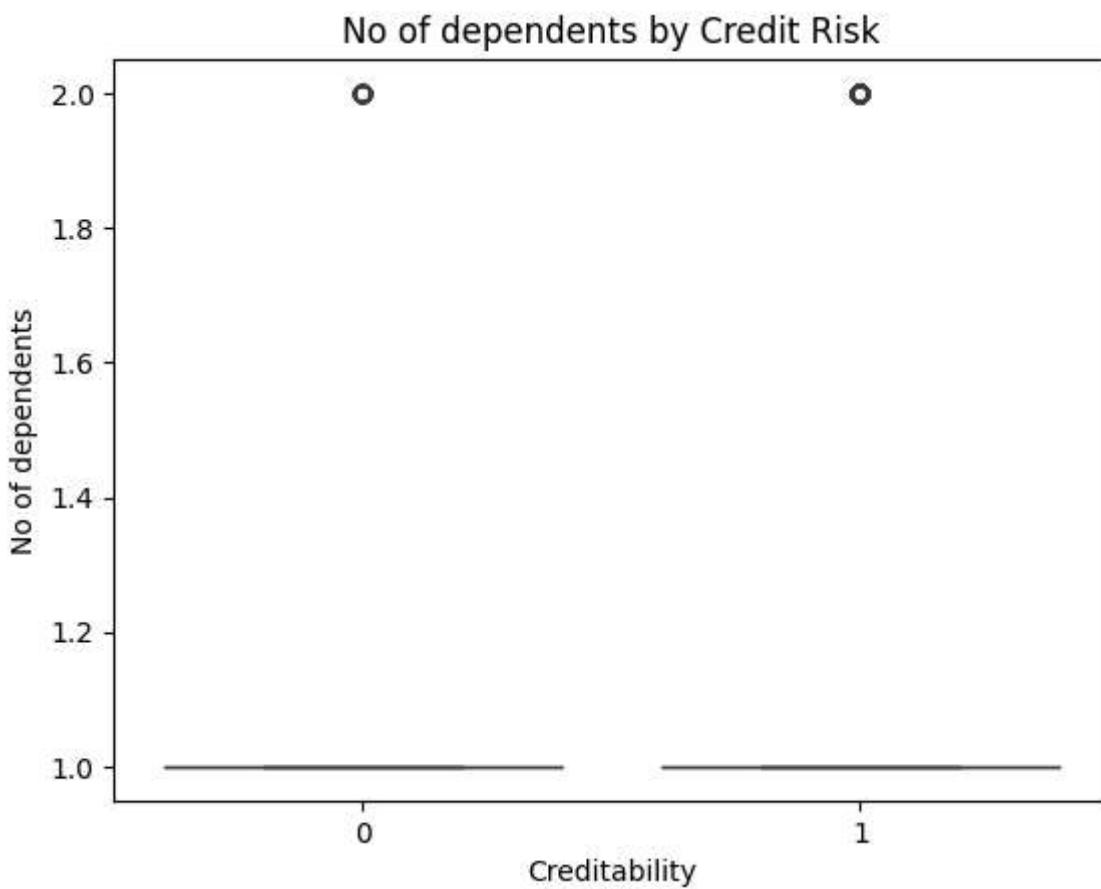
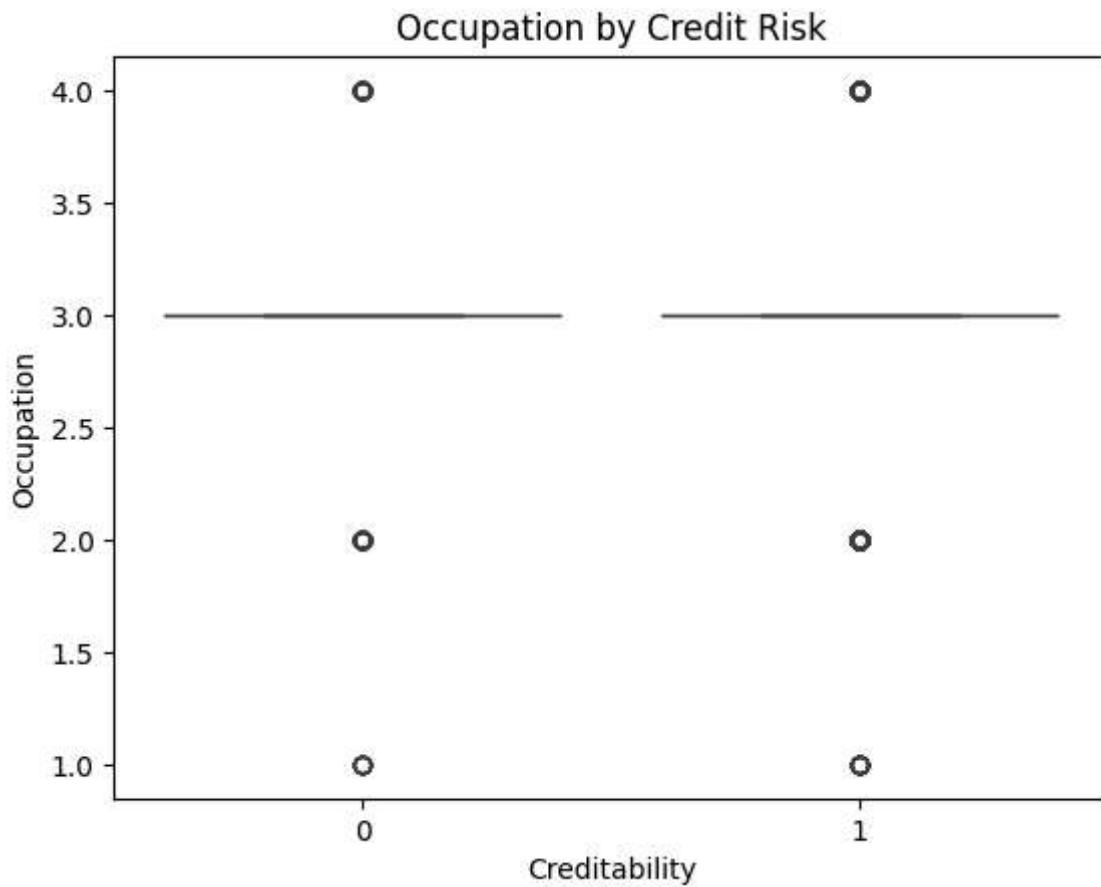


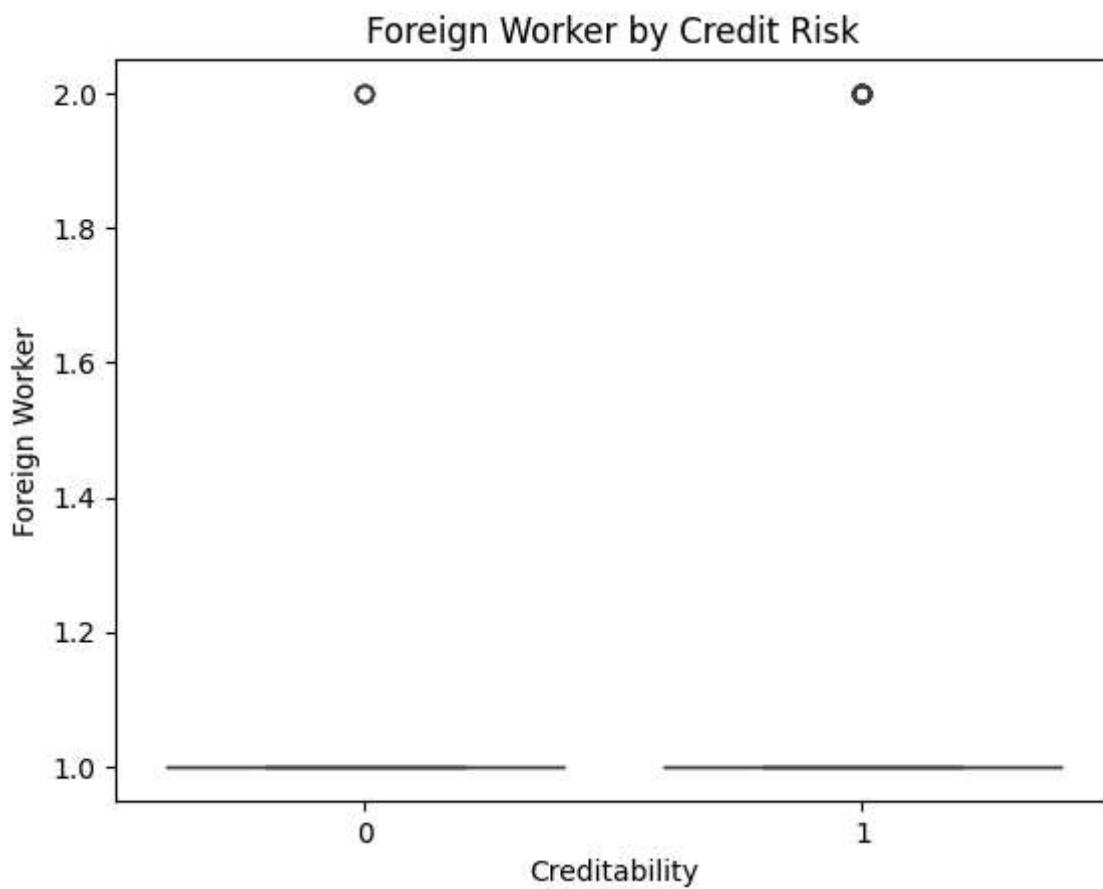
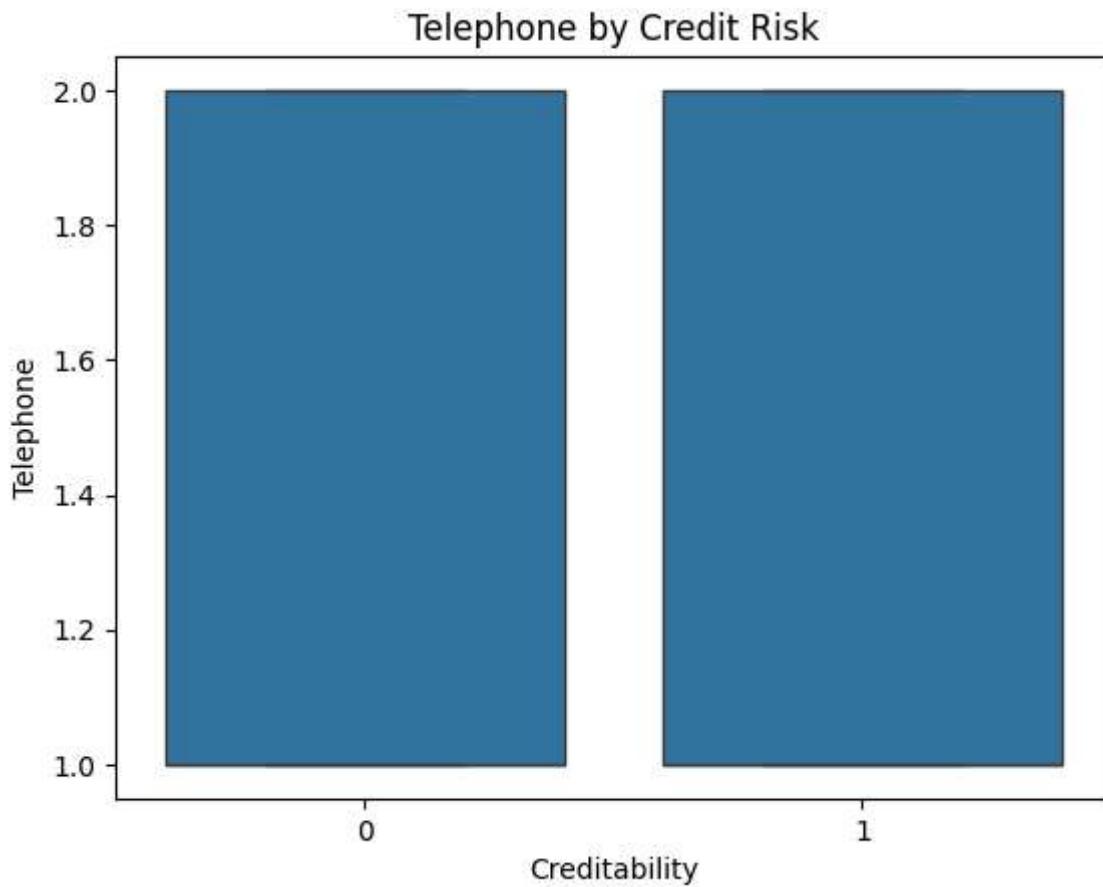
Type of apartment by Credit Risk



No of Credits at this Bank by Credit Risk



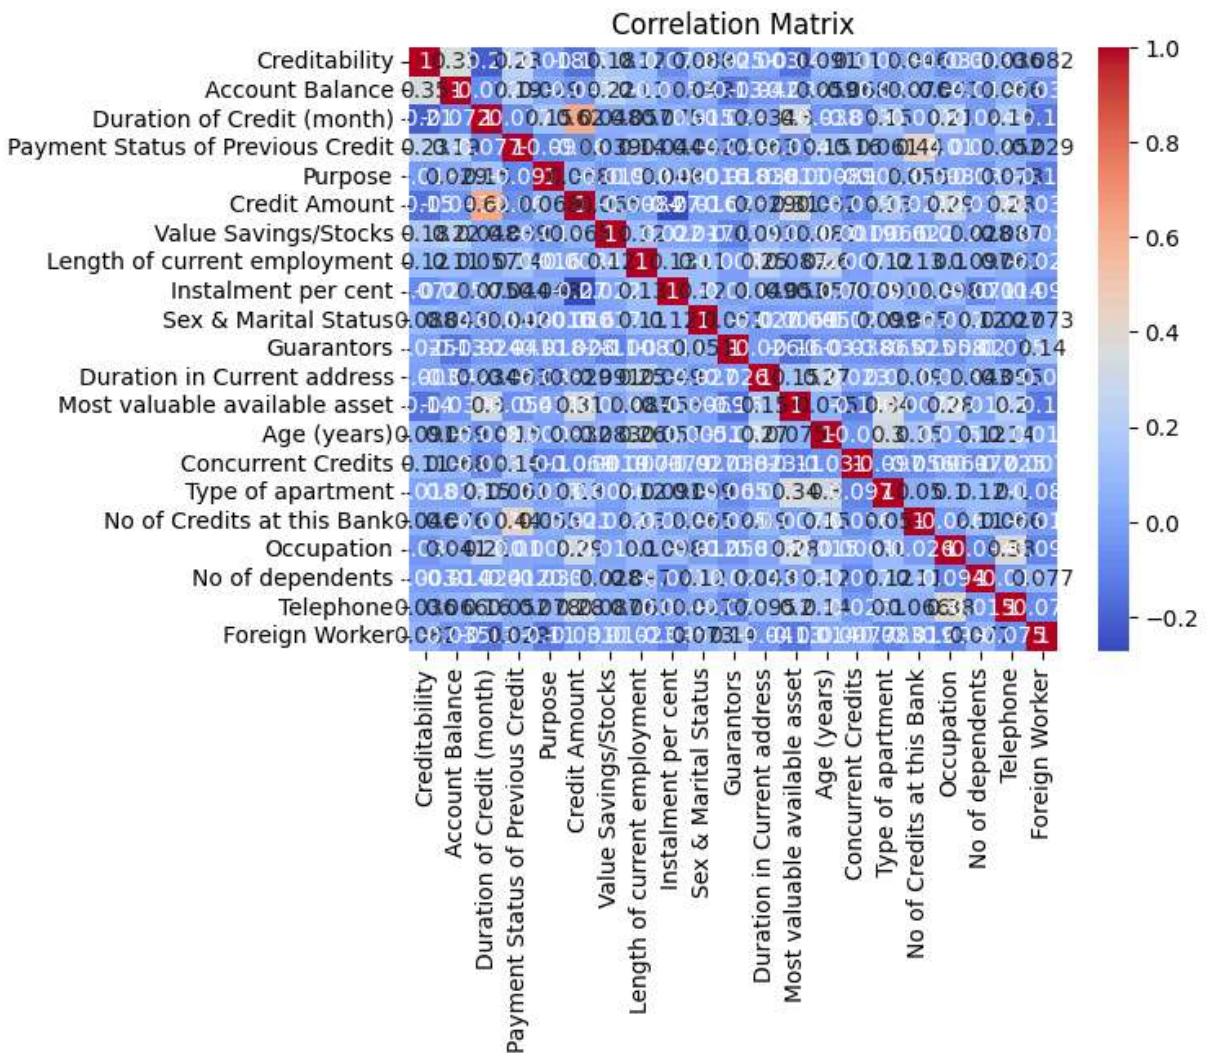




```
In [19]: for col in df.select_dtypes(include='object').columns:
    ct = pd.crosstab(df[col], df['CreditRisk'], normalize='index')
    ct.plot(kind='bar', stacked=True)
    plt.title(f"Creditability by {col}")
    plt.ylabel("Proportion")
    plt.show()
```

```
In [20]: import numpy as np

corr = df.corr(numeric_only=True)
sns.heatmap(corr, annot=True, cmap='coolwarm')
plt.title("Correlation Matrix")
plt.show()
```

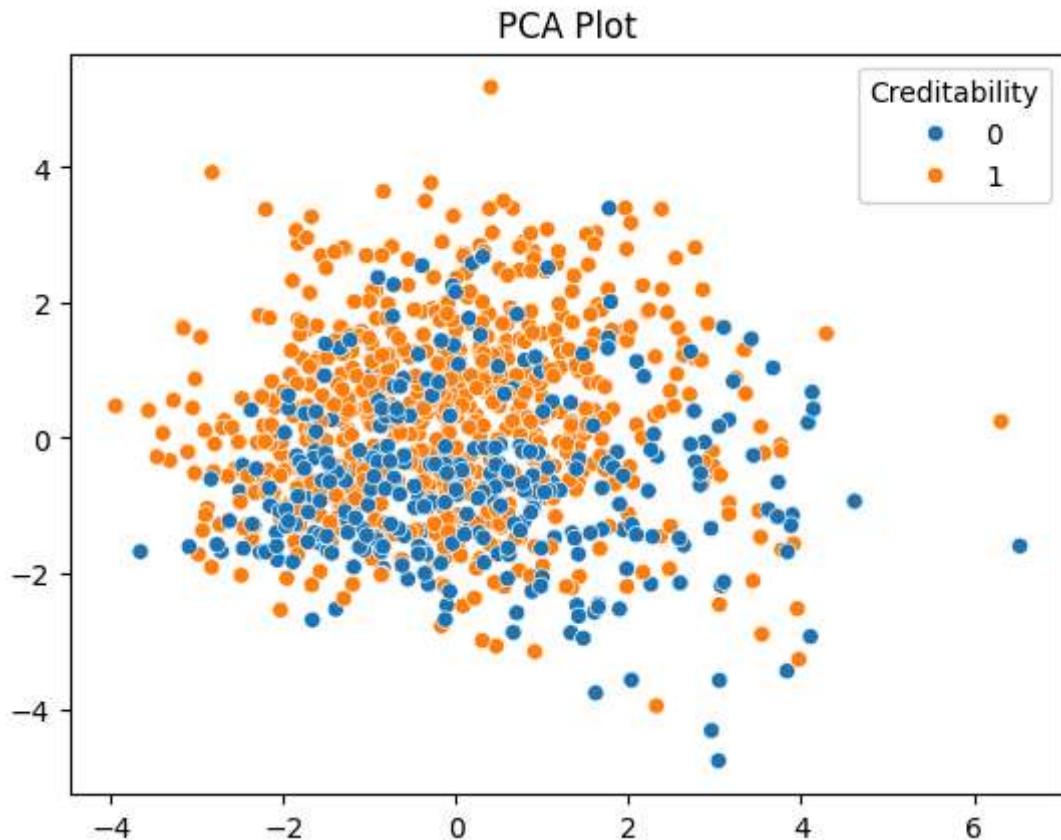


```
In [23]: from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler

X = df.drop("Creditability", axis=1).select_dtypes(include='number')
X_scaled = StandardScaler().fit_transform(X)

pca = PCA(n_components=2)
components = pca.fit_transform(X_scaled)
```

```
sns.scatterplot(x=components[:,0], y=components[:,1], hue=df['Creditability'])
plt.title("PCA Plot")
plt.show()
```



```
In [83]: print("Class balance in full target y:")
print(y.value_counts())
```

Class balance in full target y:
 Creditability
 1 1000
 Name: count, dtype: int64

```
In [85]: print(df['Creditability'].describe())
```

	Creditability
count	1000.0
mean	1.0
std	0.0
min	1.0
25%	1.0
50%	1.0
75%	1.0
max	1.0

Name: Creditability, dtype: float64

```
In [87]: import pandas as pd

df = pd.read_csv("C:/Users/Admin/Downloads/index_data.csv")
print(df['Creditability'].value_counts())
```

```
Creditability
1    700
0    300
Name: count, dtype: int64
```

```
In [88]: X = df.drop(columns=['Creditability'])
y = df['Creditability']
```

```
In [89]: X = pd.get_dummies(X, drop_first=True)
```

```
In [90]: from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

```
In [91]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(
    X_scaled, y, test_size=0.3, random_state=42, stratify=y
)
```

```
In [92]: from sklearn.linear_model import LogisticRegression
model = LogisticRegression(max_iter=1000)
model.fit(X_train, y_train)
```

Out[92]:

- ▼ LogisticRegression ⓘ ?
- ▶ Parameters

```
In [93]: from sklearn.metrics import classification_report
print(classification_report(y_test, model.predict(X_test)))
```

	precision	recall	f1-score	support
0	0.70	0.48	0.57	90
1	0.80	0.91	0.86	210
accuracy			0.78	300
macro avg	0.75	0.70	0.71	300
weighted avg	0.77	0.78	0.77	300

```
In [119...]: from sklearn.metrics import roc_curve, auc
import matplotlib.pyplot as plt

# Get probability scores for the positive class
y_score_lr = model.predict_proba(X_test)[:, 1]

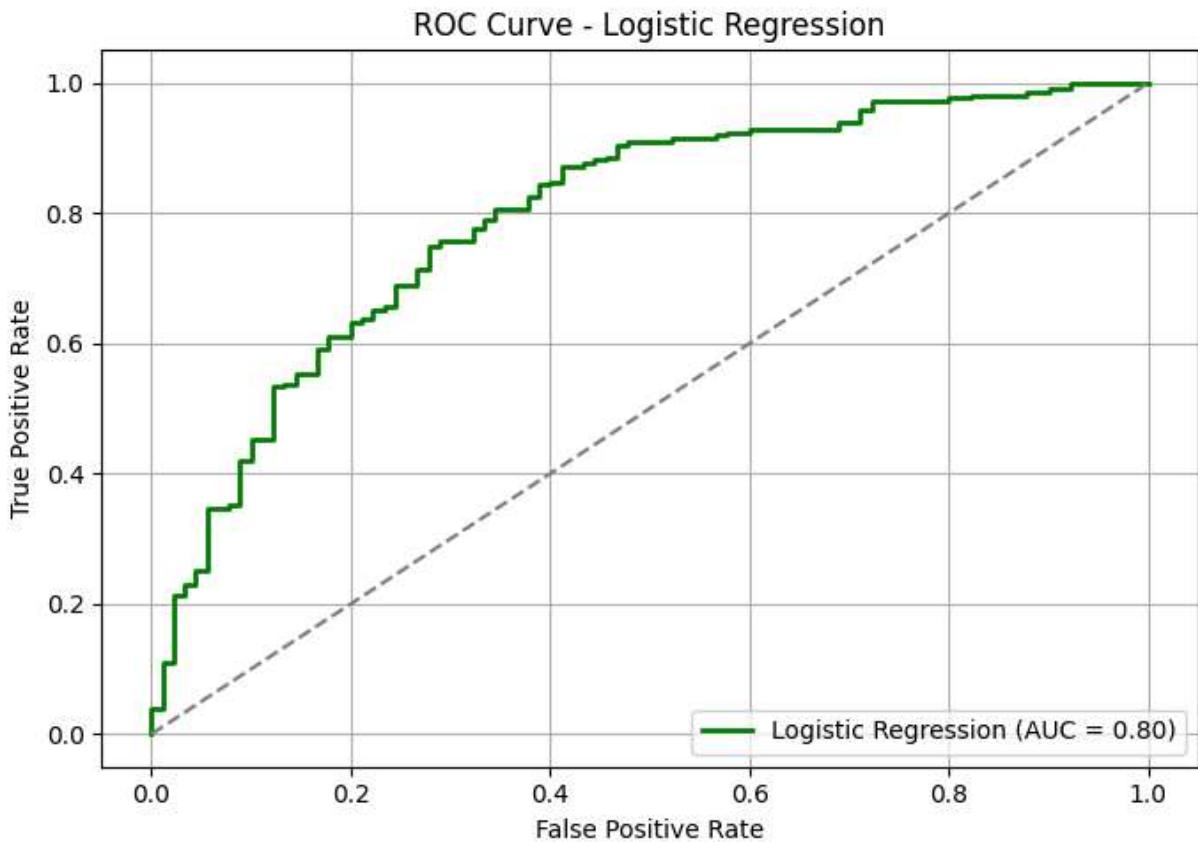
# Calculate False Positive Rate and True Positive Rate
fpr, tpr, _ = roc_curve(y_test, y_score_lr)
roc_auc = auc(fpr, tpr)

# Plot ROC curve
plt.figure(figsize=(7, 5))
plt.plot(fpr, tpr, color='green', lw=2, label=f'Logistic Regression (AUC = {roc_auc:.4f})')
plt.plot([0, 1], [0, 1], color='gray', linestyle='--')
```

```

plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve - Logistic Regression')
plt.legend(loc='lower right')
plt.grid(True)
plt.tight_layout()
plt.show()

```



In [95]:

```

import pandas as pd

# Load the correct dataset (use lrgc_data.csv if you know it's clean)
df = pd.read_csv("C:/Users/Admin/Downloads/index_data.csv")

# Confirm target balance
print(df['Creditability'].value_counts()) # Expect both 0 and 1

```

```

Creditability
1    700
0    300
Name: count, dtype: int64

```

In [96]:

```

# Separate features and target
X = df.drop(columns=['Creditability'])
y = df['Creditability']

# Encode categorical variables
X = pd.get_dummies(X, drop_first=True)

# Scale numerical features
from sklearn.preprocessing import StandardScaler

```

```
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

In [97]:

```
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(
    X_scaled, y, test_size=0.3, random_state=42, stratify=y
)
```

In [98]:

```
from sklearn.ensemble import RandomForestClassifier

rf_model = RandomForestClassifier(n_estimators=100, random_state=42)
rf_model.fit(X_train, y_train)
```

Out[98]:

▼ RandomForestClassifier ⓘ ?
► Parameters

In [99]:

```
from sklearn.metrics import classification_report, confusion_matrix

y_pred_rf = rf_model.predict(X_test)

print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred_rf))

print("\nClassification Report:")
print(classification_report(y_test, y_pred_rf))
```

Confusion Matrix:

```
[[ 40  50]
 [ 16 194]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.71	0.44	0.55	90
1	0.80	0.92	0.85	210
accuracy			0.78	300
macro avg	0.75	0.68	0.70	300
weighted avg	0.77	0.78	0.76	300

In [100...]:

```
from sklearn.metrics import roc_curve, auc
import matplotlib.pyplot as plt
```

```
# Predict probabilities
y_score_rf = rf_model.predict_proba(X_test)[:, 1]

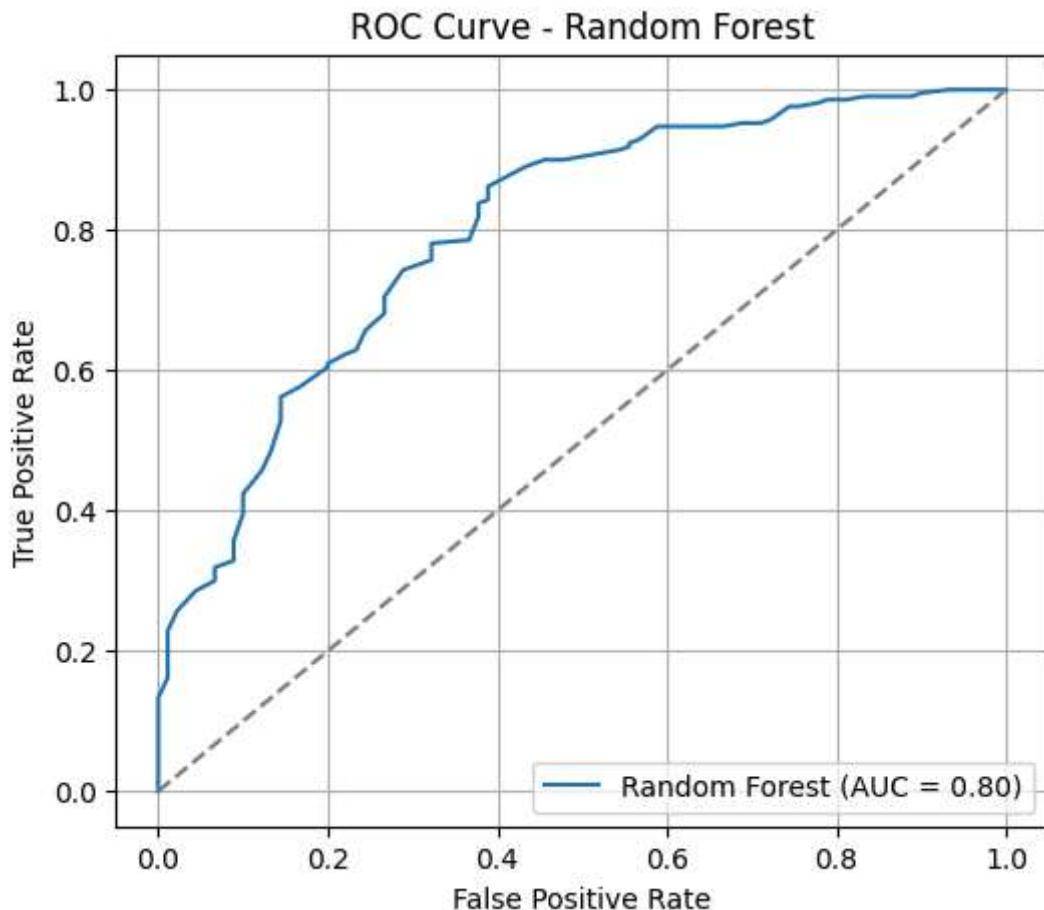
fpr, tpr, _ = roc_curve(y_test, y_score_rf)
roc_auc = auc(fpr, tpr)

plt.figure(figsize=(6, 5))
plt.plot(fpr, tpr, label=f'Random Forest (AUC = {roc_auc:.2f})')
```

```

plt.plot([0, 1], [0, 1], linestyle='--', color='gray')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve - Random Forest')
plt.legend(loc='lower right')
plt.grid(True)
plt.show()

```



In [101...]

```

import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import classification_report, confusion_matrix
from xgboost import XGBClassifier

```

In [103...]

```
df = pd.read_csv("C:/Users/Admin/Downloads/index_data.csv")
```

In [104...]

```
target_column = 'Creditability' if 'Creditability' in df.columns else df.columns[-1]
```

In [105...]

```
X = df.drop(columns=[target_column])
y = df[target_column]
```

In [106...]

```
X = pd.get_dummies(X, drop_first=True)
```

In [139...]

```
X = X.loc[y.notna()]
y = y.loc[y.notna()]
```

```
In [108...     scaler = StandardScaler()
           X_scaled = scaler.fit_transform(X)
```

```
In [109...     X_train, X_test, y_train, y_test = train_test_split(
           X_scaled, y, test_size=0.3, random_state=42, stratify=y
         )
```

```
In [110...     xgb_model = XGBClassifier(
           use_label_encoder=False,
           eval_metric='logloss',
           n_estimators=50,
           random_state=42
         )
```

```
In [138...     xgb_model.fit(X_train, y_train)
```

C:\Users\Admin\AppData\Local\Programs\Python\Python313\Lib\site-packages\xgboost\training.py:183: UserWarning: [14:23:07] WARNING: C:\actions-runner_work\xgboost\xgboost\src\learner.cc:738:
Parameters: { "use_label_encoder" } are not used.

```
bst.update(dtrain, iteration=i, fobj=obj)
```

Out[138... ▾ XGBClassifier ⓘ ⓘ
▶ Parameters

```
In [114...     y_pred = xgb_model.predict(X_test)
```

```
In [115...     print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
           print("\nClassification Report:\n", classification_report(y_test, y_pred))
```

Confusion Matrix:
[[44 46]
 [27 183]]

		precision	recall	f1-score	support
	0	0.62	0.49	0.55	90
	1	0.80	0.87	0.83	210
	accuracy			0.76	300
	macro avg	0.71	0.68	0.69	300
	weighted avg	0.75	0.76	0.75	300

```
In [116...     from sklearn.metrics import accuracy_score, classification_report

     # Predictions
     y_pred_train = model.predict(X_train)
     y_pred_test = model.predict(X_test)

     # Training performance
```

```

print("◆ Training Accuracy:", accuracy_score(y_train, y_pred_train))
print("◆ Training Report:\n", classification_report(y_train, y_pred_train))

# Testing performance
print("◆ Testing Accuracy:", accuracy_score(y_test, y_pred_test))
print("◆ Testing Report:\n", classification_report(y_test, y_pred_test))

```

◆ Training Accuracy: 0.78
◆ Training Report:

	precision	recall	f1-score	support
0	0.69	0.49	0.57	210
1	0.81	0.90	0.85	490
accuracy			0.78	700
macro avg	0.75	0.70	0.71	700
weighted avg	0.77	0.78	0.77	700

◆ Testing Accuracy: 0.7833333333333333
◆ Testing Report:

	precision	recall	f1-score	support
0	0.70	0.48	0.57	90
1	0.80	0.91	0.86	210
accuracy			0.78	300
macro avg	0.75	0.70	0.71	300
weighted avg	0.77	0.78	0.77	300

In [121...]

```

print("◆ y_train sample:")
print(y_train.head())

print("\n◆ y_test sample:")
print(y_test.head())

```

◆ y_train sample:

247	1
667	1
191	1
181	1
368	1

Name: Creditability, dtype: int64

◆ y_test sample:

521	0
124	1
51	1
366	1
602	1

Name: Creditability, dtype: int64

```
In [122... print("◆ y_train class distribution:")
print(y_train.value_counts())

print("\n◆ y_test class distribution:")
print(y_test.value_counts())

◆ y_train class distribution:
Creditability
1    490
0    210
Name: count, dtype: int64

◆ y_test class distribution:
Creditability
1    210
0     90
Name: count, dtype: int64
```

```
In [123... print("y_train shape:", y_train.shape)
print("y_test shape:", y_test.shape)

y_train shape: (700,)
y_test shape: (300,)
```

```
In [124... print("y_train class distribution:")
print(y_train.value_counts())
print("\ny_test class distribution:")
print(y_test.value_counts())

y_train class distribution:
Creditability
1    490
0    210
Name: count, dtype: int64

y_test class distribution:
Creditability
1    210
0     90
Name: count, dtype: int64
```

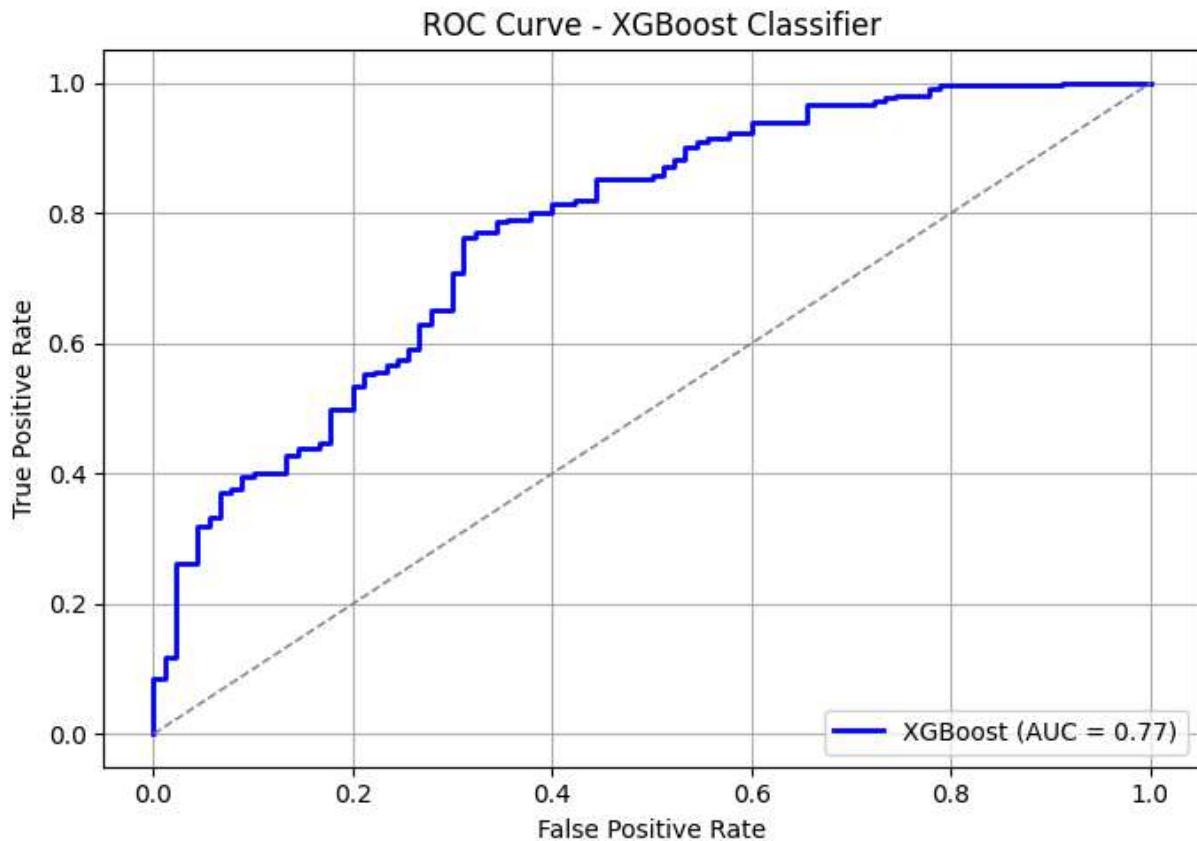
```
In [117... from sklearn.metrics import roc_curve, auc
import matplotlib.pyplot as plt

# Predict probabilities for test set
y_score = xgb_model.predict_proba(X_test)[:, 1]

# Compute ROC curve and AUC
fpr, tpr, _ = roc_curve(y_test, y_score)
roc_auc = auc(fpr, tpr)

# Plot ROC curve
plt.figure(figsize=(7, 5))
plt.plot(fpr, tpr, color='blue', lw=2, label=f'XGBoost (AUC = {roc_auc:.2f})')
plt.plot([0, 1], [0, 1], color='gray', lw=1, linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
```

```
plt.title('ROC Curve - XGBoost Classifier')
plt.legend(loc='lower right')
plt.grid(True)
plt.tight_layout()
plt.show()
```



In [135...]:

```
from sklearn.metrics import confusion_matrix

print("Logistic Regression")
print(confusion_matrix(y_test, y_pred_lr))

print("\nRandom Forest")
print(confusion_matrix(y_test, y_pred_rf))

print("\nXGBoost")
print(confusion_matrix(y_test, y_pred_xgb))
```

Logistic Regression

```
[[ 43  47]
 [ 18 192]]
```

Random Forest

```
[[ 40  50]
 [ 16 194]]
```

XGBoost

```
[[ 43  47]
 [ 28 182]]
```

```
In [126...]: from sklearn.linear_model import LogisticRegression
lr_model = LogisticRegression(max_iter=1000)
lr_model.fit(X_train, y_train)
```

Out[126...]:

LogisticRegression		
▼ Parameters		
penalty	'l2'	
dual	False	
tol	0.0001	
C	1.0	
fit_intercept	True	
intercept_scaling	1	
class_weight	None	
random_state	None	
solver	'lbfgs'	
max_iter	1000	
multi_class	'deprecated'	
verbose	0	
warm_start	False	
n_jobs	None	
l1_ratio	None	

```
In [134...]: y_pred_lr = lr_model.predict(X_test)
```

```
In [133...]: y_pred_xgb = xgb_model.predict(X_test)
```

```
In [136...]: from sklearn.metrics import classification_report
```

```
print("Logistic Regression:\n", classification_report(y_test, y_pred_lr))
print("Random Forest:\n", classification_report(y_test, y_pred_rf))
print("XGBoost:\n", classification_report(y_test, y_pred_xgb))
```

Logistic Regression:

	precision	recall	f1-score	support
0	0.70	0.48	0.57	90
1	0.80	0.91	0.86	210
accuracy			0.78	300
macro avg	0.75	0.70	0.71	300
weighted avg	0.77	0.78	0.77	300

Random Forest:

	precision	recall	f1-score	support
0	0.71	0.44	0.55	90
1	0.80	0.92	0.85	210
accuracy			0.78	300
macro avg	0.75	0.68	0.70	300
weighted avg	0.77	0.78	0.76	300

XGBoost:

	precision	recall	f1-score	support
0	0.61	0.48	0.53	90
1	0.79	0.87	0.83	210
accuracy			0.75	300
macro avg	0.70	0.67	0.68	300
weighted avg	0.74	0.75	0.74	300

In [137]:

```

from sklearn.metrics import roc_curve, auc
import matplotlib.pyplot as plt

# Get probabilities
lr_probs = lr_model.predict_proba(X_test)[:, 1]
rf_probs = rf_model.predict_proba(X_test)[:, 1]
xgb_probs = xgb_model.predict_proba(X_test)[:, 1]

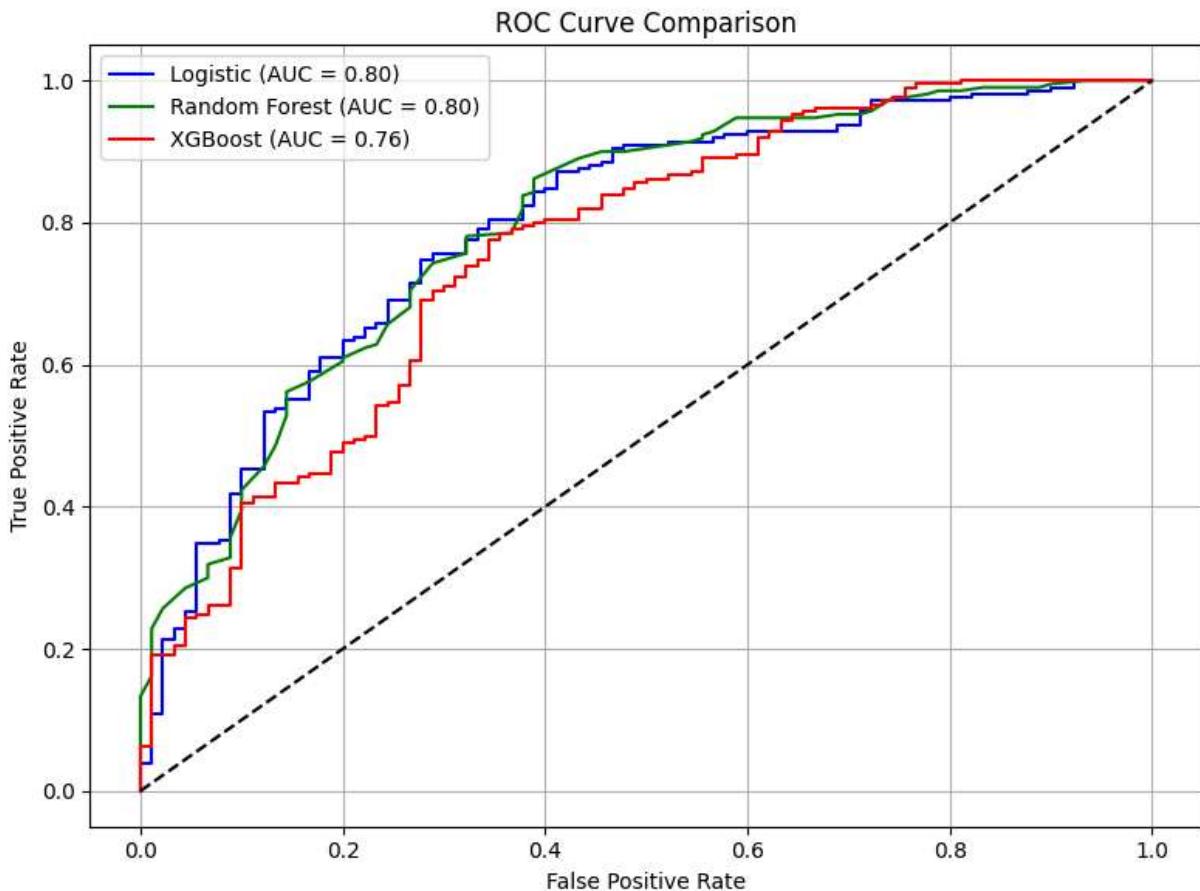
# Calculate ROC curve and AUC
fpr_lr, tpr_lr, _ = roc_curve(y_test, lr_probs)
fpr_rf, tpr_rf, _ = roc_curve(y_test, rf_probs)
fpr_xgb, tpr_xgb, _ = roc_curve(y_test, xgb_probs)

auc_lr = auc(fpr_lr, tpr_lr)
auc_rf = auc(fpr_rf, tpr_rf)
auc_xgb = auc(fpr_xgb, tpr_xgb)

# Plot
plt.figure(figsize=(8, 6))
plt.plot(fpr_lr, tpr_lr, label=f"Logistic (AUC = {auc_lr:.2f})", color='blue')
plt.plot(fpr_rf, tpr_rf, label=f"Random Forest (AUC = {auc_rf:.2f})", color='green')
plt.plot(fpr_xgb, tpr_xgb, label=f"XGBoost (AUC = {auc_xgb:.2f})", color='red')
plt.plot([0, 1], [0, 1], 'k--')
plt.title('ROC Curve Comparison')
plt.xlabel('False Positive Rate')

```

```
plt.ylabel('True Positive Rate')
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()
```



In [140]:

```
print("X_train shape:", X_train.shape)
print("X_test shape:", X_test.shape)
print("y_train shape:", y_train.shape)
print("y_test shape:", y_test.shape)
```

```
X_train shape: (700, 20)
X_test shape: (300, 20)
y_train shape: (700,)
y_test shape: (300,)
```

In [141]:

```
print("y_train class distribution:\n", y_train.value_counts(normalize=True))
print("y_test class distribution:\n", y_test.value_counts(normalize=True))
```

```
y_train class distribution:
 Creditability
 1    0.7
 0    0.3
 Name: proportion, dtype: float64
y_test class distribution:
 Creditability
 1    0.7
 0    0.3
 Name: proportion, dtype: float64
```

```
In [142...  
import pandas as pd  
X_train_df = pd.DataFrame(X_train)  
X_test_df = pd.DataFrame(X_test)  
  
print("\nX_train summary:\n", X_train_df.describe())  
print("\nX_test summary:\n", X_test_df.describe())
```

X_train summary:

	0	1	2	3	4	5	\
count	700.000000	700.000000	700.000000	700.000000	700.000000	700.000000	
mean	-0.011251	-0.003212	-0.020454	-0.024790	-0.035779	-0.012212	
std	1.007111	0.998880	0.993524	0.994291	0.969851	0.997579	
min	-1.254566	-1.402415	-2.350870	-1.030963	-1.070856	-0.699707	
25%	-1.254566	-0.738668	-0.503428	-0.666407	-0.683894	-0.699707	
50%	-0.459026	-0.240857	-0.503428	-0.301852	-0.360909	-0.699707	
75%	1.132053	0.256953	1.344014	0.062704	0.210539	0.566731	
max	1.132053	3.243815	1.344014	2.614592	4.492104	1.833169	
	6	7	8	9	10	11	\
count	700.000000	700.000000	700.000000	700.000000	700.000000	700.000000	
mean	-0.009226	-0.021847	-0.016956	0.037400	0.011007	0.003266	
std	1.009840	1.006961	0.992341	1.049240	1.012296	0.996993	
min	-1.973997	-1.764514	-2.376626	-0.303686	-1.672459	-1.293723	
25%	-0.317959	-0.870183	-0.963650	-0.303686	-0.765977	-1.293723	
50%	-0.317959	0.024147	0.449326	-0.303686	0.140505	-0.341055	
75%	1.338078	0.918477	0.449326	-0.303686	1.046987	0.611613	
max	1.338078	0.918477	1.862303	3.885083	1.046987	1.564281	
	12	13	14	15	16	17	\
count	700.000000	700.000000	700.000000	700.000000	700.000000	700.000000	
mean	0.013924	-0.011141	0.017253	-0.017073	-0.001749	0.005921	
std	1.016168	1.008277	1.002401	0.988305	1.005136	1.006330	
min	-1.457831	-2.375050	-1.751205	-0.704926	-2.914492	-0.428290	
25%	-0.752799	0.460831	0.135869	-0.704926	0.146949	-0.428290	
50%	-0.224024	0.460831	0.135869	-0.704926	0.146949	-0.428290	
75%	0.569138	0.460831	0.135869	1.027079	0.146949	-0.428290	
max	3.477397	0.460831	2.022944	4.491089	1.677670	2.334869	
	18	19					
count	700.000000	700.000000					
mean	-0.011063	0.038597					
std	0.998485	1.090667					
min	-0.823318	-0.196014					
25%	-0.823318	-0.196014					
50%	-0.823318	-0.196014					
75%	1.214598	-0.196014					
max	1.214598	5.101669					

X_test summary:

	0	1	2	3	4	5	\
count	300.000000	300.000000	300.000000	300.000000	300.000000	300.000000	
mean	0.026253	0.007495	0.047726	0.057843	0.083485	0.028495	
std	0.986069	1.005904	1.016664	1.014156	1.065616	1.008381	
min	-1.254566	-1.402415	-2.350870	-1.030963	-1.061640	-0.699707	
25%	-0.657911	-0.738668	-0.503428	-0.666407	-0.649159	-0.699707	
50%	-0.459026	-0.240857	-0.503428	0.062704	-0.297110	-0.699707	
75%	1.132053	0.568085	1.344014	0.062704	0.399634	0.566731	
max	1.132053	4.239436	1.344014	2.614592	5.370764	1.833169	
	6	7	8	9	10	11	\
count	300.000000	300.000000	300.000000	300.000000	300.000000	300.000000	
mean	0.021528	0.050977	0.039563	-0.087266	-0.025684	-0.007621	
std	0.979690	0.985040	1.019862	0.871709	0.973600	1.010269	

min	-1.973997	-1.764514	-2.376626	-0.303686	-1.672459	-1.293723
25%	-0.317959	-0.870183	-0.963650	-0.303686	-0.765977	-1.293723
50%	-0.317959	0.024147	0.449326	-0.303686	0.140505	-0.341055
75%	1.338078	0.918477	0.449326	-0.303686	1.046987	0.611613
max	1.338078	0.918477	1.862303	3.885083	1.046987	1.564281

	12	13	14	15	16	17	\
count	300.000000	300.000000	300.000000	300.000000	300.000000	300.000000	
mean	-0.032490	0.025996	-0.040258	0.039836	0.004082	-0.013816	
std	0.963830	0.983290	0.996556	1.028967	0.991256	0.988293	
min	-1.369702	-2.375050	-1.751205	-0.704926	-2.914492	-0.428290	
25%	-0.752799	0.460831	0.135869	-0.704926	0.146949	-0.428290	
50%	-0.224024	0.460831	0.135869	-0.704926	0.146949	-0.428290	
75%	0.481009	0.460831	0.135869	1.027079	0.146949	-0.428290	
max	3.477397	0.460831	2.022944	4.491089	1.677670	2.334869	

	18	19
count	300.000000	300.000000
mean	0.025814	-0.090061
std	1.006385	0.742915
min	-0.823318	-0.196014
25%	-0.823318	-0.196014
50%	-0.823318	-0.196014
75%	1.214598	-0.196014
max	1.214598	5.101669

In [143]:

```
from sklearn.metrics import accuracy_score

train_preds = model.predict(X_train)
test_preds = model.predict(X_test)

print("Train Accuracy:", accuracy_score(y_train, train_preds))
print("Test Accuracy :", accuracy_score(y_test, test_preds))
```

Train Accuracy: 0.78

Test Accuracy : 0.7833333333333333

In []: